



Uniwersytet Ekonomiczny  
we Wrocławiu

Program

**Business informatics**

**Dawid Jeleńkowski**

Student No. 174682

**MASTER'S THESIS**

# **Leveraging Functional Programming in Scala for Efficient Data Engineering**

**Wykorzystanie programowania funkcyjnego w Scala do wydajnej inżynierii danych**

Master's thesis written under the supervision of

**dr inż. Adam Sulich, prof. UEW**

I approve the thesis and I request for further processing

---

Supervisor's signature

WROCLAW 2024

## **Abstract**

The potential of functional programming in Scala for efficient data engineering is explored in this thesis. As organizations grapple with ever-increasing volumes of data, the need for scalable, maintainable, and performant data processing systems has become paramount. Scala, a multi-paradigm language that seamlessly blends object-oriented and functional programming concepts, is examined as a powerful tool for addressing these challenges. The core principles of functional programming in Scala are investigated, including immutable data structures, higher-order functions, lazy evaluation, pattern matching, algebraic data types, type classes, and monads. These concepts are demonstrated to provide a solid foundation for building robust data pipelines. The entire data engineering workflow is analyzed, from data ingestion to serving, with a focus on how functional programming paradigms can be applied at each stage. Parallel and distributed data processing techniques are explored, showcasing Scala's capabilities in handling large-scale data operations. Particular attention is given to stream processing, reflecting its growing importance in modern data architectures. Popular stream processing frameworks in Scala are evaluated, and their strengths in real-time data analysis are highlighted. The critical aspects of testing, deploying, and monitoring data pipelines are addressed. Various testing methodologies tailored for Scala are examined, underlining the importance of comprehensive testing in ensuring reliable and efficient data processing systems. Deployment strategies, including cloud deployment, are discussed, along with best practices for monitoring and optimizing data pipelines in production environments. Through this comprehensive exploration, it is demonstrated that functional programming in Scala offers a powerful toolkit for addressing the complexities of modern data engineering.

## Abstract

W niniejszej rozprawie zbadano potencjał programowania funkcjonalnego w języku Scala w zakresie wydajnej inżynierii danych. W miarę jak organizacje zmagają się z coraz większymi ilościami danych, potrzeba skalowalnych, łatwych w utrzymaniu i wydajnych systemów przetwarzania danych stała się nadrzędna. Scala, wieloparadygmataowy język, który płynnie łączy koncepcje programowania obiektowego i funkcjonalnego, jest badany jako potężne narzędzie do sprostania tym wyzwaniom. Badane są podstawowe zasady programowania funkcyjnego w Scali, w tym immutable data structures, higher-order functions, lazy evaluation, pattern matching, algebraic data types, type classes i monady. Koncepcje te są demonstrowane w celu zapewnienia solidnych podstaw do budowania wydajnych strumieni danych. Analizowany jest cały cykl pracy związany z inżynierią danych, od pozyskiwania danych po ich udostępnianie, z naciskiem na to, w jaki sposób paradygmaty programowania funkcjonalnego mogą być stosowane na każdym etapie. Badane są techniki równoległego i rozproszonego przetwarzania danych, prezentując możliwości Scali w zakresie obsługi operacji na danych na dużą skalę. Szczególną uwagę poświęcono przetwarzaniu strumieniowemu, odzwierciedlając jego rosnące znaczenie w nowoczesnych architekturach danych. Oceniono popularne biblioteki przetwarzania strumieniowego w Scali i podkreślono ich mocne strony w analizie danych w czasie rzeczywistym. Omówiono krytyczne aspekty testowania, wdrażania i monitorowania strumieni danych. Przeanalizowano różne metodologie testowania dostosowane do Scali, podkreślając znaczenie kompleksowego testowania w zapewnianiu niezawodnych i wydajnych systemów przetwarzania danych. Omówiono strategie wdrażania, w tym wdrażanie w chmurze, wraz z najlepszymi praktykami monitorowania i optymalizacji potoków danych w środowiskach produkcyjnych. Dzięki tej wszechstronnej eksploracji wykazano, że programowanie funkcjonalne w Scali oferuje potężny zestaw narzędzi do radzenia sobie ze złożonością nowoczesnej inżynierii danych.

# Contents

<b>Introduction</b>	<b>1</b>
Background and rationale . . . . .	1
Research objectives and learning goals . . . . .	1
Purpose and scope . . . . .	2
Research methods . . . . .	2
Thesis structure . . . . .	2
 <b>1 Core Principles of Functional Programming in Scala for Data Engineering</b>	 <b>5</b>
1.1 Background . . . . .	6
1.2 Growing a Language . . . . .	6
1.3 What Makes Scala Scalable? . . . . .	7
1.4 Object-Oriented Paradigm in Scala . . . . .	8
1.5 Overview of Functional Programming Paradigm . . . . .	9
1.6 Core Features of Functional Programing . . . . .	9
1.6.1 Immutable data structures . . . . .	9
1.6.2 Higher-Order Functions . . . . .	11
1.6.3 Lazy Evaluation . . . . .	13
1.6.4 Pattern Matching and Algebraic Data Types (ADTs) . . . . .	16

1.6.5	Type Classes	18
1.6.6	Monads and error handling	19
1.6.7	Parallel and Distributed Processing	21
<b>2</b>	<b>Data Engineering Workflow</b>	<b>23</b>
2.1	Data Quality Dimensions	24
2.2	Data Ingestion	25
2.3	Data Processing	25
2.4	Data Storage Systems	26
2.4.1	The Role and Challenges Associated with Storing	27
2.4.2	Types of Data Storage Systems	28
2.5	Data Integration	30
2.5.1	Data Pipeline Architectures	30
2.5.2	ETL (Extract, Transform, Load)	31
2.5.3	ELT (Extract, Load, Transform)	32
2.5.4	Comparison of ETL and ELT	32
2.6	Data Security and Governance	33
2.7	Real-World Data Security Challenges and Solutions	34
2.8	Data Serving	36
2.9	Monitoring and Maintenance	37
<b>3</b>	<b>Parallel and distributed data processing</b>	<b>38</b>
3.1	Fundamentals of Parallel and Distributed Processing for Big Data	39
3.2	Distributed Processing Frameworks for Big Data	40
3.3	Architectural Patterns for Big Data Processing	40

3.4	Parallel Collections	43
3.4.1	Task Support (TS)	44
3.4.2	Performance Implications of Task Support	45
3.5	Futures and Async Programming	46
3.6	Akka Framework	46
3.6.1	Actor Model	47
3.6.2	Actors Advantage in Data Engineering	49
3.6.3	Reactive Streams	50
3.7	Apache Spark	54
3.7.1	Spark DataFrames and APIs	55
<b>4</b>	<b>Stream processing with Scala</b>	<b>57</b>
4.1	Definition and Concepts	58
4.2	Introduction to Stream Processing	58
4.2.1	Evolution of Data Processing Paradigms	58
4.3	Stream Processing in Modern Data Engineering	59
4.4	Stream Processing in Big Data Ecosystems	61
4.5	Challenges and Considerations in Stream Processing	62
4.6	Stream Processing Technologies and Frameworks	64
4.7	Popular Stream Processing Frameworks in Scala	64
<b>5</b>	<b>Testing Data Pipelines</b>	<b>66</b>
5.1	Testing Levels for Data Pipelines	67
5.2	Testing Frameworks and Tools for Scala	67
<b>6</b>	<b>Deploying and monitoring data pipelines</b>	<b>69</b>

6.1 Cloud Deployment . . . . .	70
<b>7 Conclusion</b>	<b>72</b>
<b>Bibliography</b>	<b>74</b>

# Introduction

## Background and rationale

Data-driven society of today means that companies rely more and more on insights from data sets to make wise choices, optimize procedures, and keep one step ahead of the competition. The need of data engineers in building and overseeing the physical infrastructure and data pipelines has grown exponentially with the amount of data.

In this sense, building data pipelines has come to be a paradigm for functional programming. Because it is so adaptable and works well with the Java environment, Scala, a programming language that blends object-oriented and functional programming concepts, has drawn interest in the data engineering domain.

The choice to concentrate on this subject comes from the growing importance of data engineering in the big data age and the need for data engineers to adopt methods and technologies that make it possible to build scalable pipelines. Within Scala, functional programming offers a solution to the problems data engineers have when effectively processing huge datasets.

This work is both theoretical and empirical as it shows how to use Scala's functional programming to efficiently design data while also providing a theoretical discourse. My master's thesis subject was further motivated by the way theory and practice were integrated.

## Research objectives

The primary goal of this thesis is to investigate the use of programming, in Scala, for creating scalable and maintainable data pipelines. Emphasizing understanding the principles and techniques. The research was driven by a desire to gain practical knowledge and skills in data engineering using Scala.



## Purpose and scope of the study

The goal of this thesis is to investigate the use of programming, in Scala, for constructing scalable data pipelines. The study examines the principles, libraries, and frameworks within the Scala ecosystem that are relevant to data engineering.

The scope of the work covers the data engineering workflow, from data acquisition and storage to, processing, delivery, and monitoring. In addition, it considers employing the Scala for both batch and streaming data pipelines, deliberating on the advantages and best practices associated with each method.

## Research methods

This thesis employs a combination of theoretical analysis and practical exploration to investigate the use of functional programming in Scala for data engineering. An extensive **literature review** on functional programming concepts, Scala language features, and data engineering principles was conducted to establish the theoretical foundation. Key concepts in functional programming and data engineering were analyzed and synthesized to understand their applicability and benefits in building data pipelines. Practical code examples in Scala were developed and presented throughout the thesis to illustrate how functional programming concepts can be applied to data engineering tasks. Different approaches, frameworks, and libraries within the Scala ecosystem for data processing were examined and compared to evaluate their strengths and use cases. Hypothetical scenarios and real-world inspired examples were used to demonstrate the application of Scala and functional programming principles to data engineering problems. Various data pipeline architectures and processing paradigms were analyzed to understand how functional programming in Scala can be leveraged in different contexts.

## Thesis structure

Chapter 1 "Core Principles of Functional Programming in Scala for Data Engineering" explores how Scala, can be leveraged for data engineering. It covers:

- Background on Scala's development and design philosophy;
- Core features of functional programming in Scala:
  - Immutable data structures,
  - Higher-order functions,
  - Lazy evaluation,
  - Pattern matching and algebraic data types,

- Type classes,
- Monads and error handling,
- Parallel and distributed processing capabilities;
- The chapter provides code examples to illustrate how these concepts can be applied to data engineering scenarios.

Chapter 2 "Data Engineering Workflow" provides a detailed overview of the data engineering workflow, covering:

- Data quality dimensions (accuracy, completeness, consistency, timeliness, uniqueness),
- Data ingestion processes,
- Data processing techniques,
- Data storage systems (data warehouses, data lakes, NoSQL databases),
- Data integration and pipeline architectures (ETL vs ELT),
- Data security and governance,
- Data serving,
- Monitoring and maintenance of data pipelines.

Chapter 3 "Parallel and Distributed Data Processing" explores how Scala's features can be utilized for efficient parallel and distributed data processing; Focuses on techniques for processing large-scale data, including:

- Fundamentals of parallel and distributed processing for big data,
- Distributed processing frameworks (e.g., Apache Hadoop, Apache Spark),
- Architectural patterns for big data processing (e.g., Lambda and Kappa architectures),
- Parallel collections in Scala,
- Futures and asynchronous programming,
- The Akka framework and actor model,
- Apache Spark's DataFrames and APIs.

Chapter 4 "Stream Processing with Scala" discusses how stream processing fits into the broader data engineering landscape and its implementation using Scala-based tools; Covers stream processing concepts and their implementation in Scala:

- Core concepts of stream processing,
- Evolution of data processing paradigms,
- Stream processing in modern data engineering and big data ecosystems,
- Challenges and considerations in stream processing,
- Stream processing technologies and frameworks,,
- Popular stream processing frameworks in Scala (e.g., Apache Spark Structured Streaming, Akka Streams).

Chapter 5 "Testing Data Pipelines" emphasizes the importance of comprehensive testing in ensuring reliable and efficient data processing systems; Focuses on testing methodologies for data pipelines:

- Testing levels for data pipelines (unit, integration, end-to-end);
- Testing frameworks and tools specific to Scala;
- Approaches for testing data quality, performance, and scalability;
- Challenges in testing data pipelines and strategies to address them.

Chapter 6 "Deploying and Monitoring Data Pipelines" discusses best practices for maintaining and optimizing data pipelines in production environments; Covers the operational aspects of data pipelines:

- Deployment strategies, including cloud deployment;
- Monitoring techniques and tools for data pipelines;
- Performance optimization and troubleshooting;
- Scaling considerations for data pipelines;
- Continuous integration and deployment (CI/CD) for data engineering projects.

# Chapter 1

## Core Principles of Functional Programming in Scala for Data Engineering

In recent years, the field of information technology has experienced exponential growth. To address the demands of managing and analyzing ever-growing data, certain solutions were created. As organizations grapple with the challenges of big data, they require data processing systems that can be scaled, efficient, and easy to maintain. Functional programming has emerged as a potent strategy for constructing such systems, delivering advantages like immutability, modularity, and parallelism that align with data-driven demands.

The chapter explores how Scala, a multi-paradigm language that seamlessly blends object-oriented and functional programming, can be leveraged to create robust data engineering solutions. The core principles of functional programming in Scala and how they are applied specifically to data engineering tasks were examined.

The chapter begins by discussing the history and design philosophy of Scala, highlighting its ability to "grow" as a language through the incorporation of domain-specific languages. Afterward, acquire knowledge regarding the critical components that render Scala an effective tool for data engineering, including its compatibility with both functional and object-oriented approaches.

1. **Immutable data structures (IDS)** and their benefits for data consistency and concurrency;
2. **Higher-order functions (HOFs)** for building flexible and reusable data processing pipelines;
3. **Lazy evaluation (LE)** for efficient handling of large datasets;
4. **Pattern matching (PM)** and **algebraic data types (ADTs)** for expressive data modeling and transformation;
5. **Type classes (TC)** for creating generic, extensible abstractions;

6. **Monads (M)** and functional error handling for robust data pipelines;
7. **Parallel and distributed processing (PDP)** capabilities.

Throughout the chapter, code examples were provided to illustrate how these concepts can be implemented in practical data engineering scenarios.

## 1.1 Background

Scala stands for a “**scalable language**”. Scala was created by German computer scientist **Martin Odersky**. Odersky worked on the Java compiler and generics development at Sun Microsystems. He wanted to make a language that combines both object-oriented and functional programming styles. He wanted to make a typed language that works with the **Java Virtual Machine (JVM)**. The very first version of Scala appeared in 2003, marking the beginning of its journey. Over time, it has become a widely used language with a plethora of libraries and tools (Odersky et al., 2006)<sup>1</sup>.

Scala is a language that combines ideas from **object-oriented and functional programming**. The design is inspired by several other languages. Scala uses **Java’s syntax and object-oriented features**, so it’s easy for Java developers to learn and works well with Java code. Scala has many things in common with Haskell and ML, like having immutable data structures, higher-order functions and pattern matching. Scala uses the ML family of languages for its type system. It has a static type system with a type interface. Scala also tries to be as short and simple as scripting languages like Python. It has features like operator overloading and implicit conversions that make DSLs (Odersky et al., 2021)<sup>2</sup>.

## 1.2 Growing a Language

Scala’s was designed with the idea of “**growing a language**”, which sets it apart from other programming languages. The language was designed to give the user only the core functionalities and the ability to build libraries upon them. Meaning, user on its own designs a new language based on Scala for solving the particular problem or issue (Odersky et al., 2021)<sup>2</sup>.

This concept is called **domain-specific languages** or **DSLs**, for short. In practice, these are **mini-languages** or application programming interfaces (**APIs**). These languages can be built in an expressive and intuitive manner for the end user. Meaning that they can be read and understood by people who have nothing to do with programming

---

<sup>1</sup>Odersky, M., Altherr, P., Cremet, V., Dragos, I., Dubochet, G., Emir, B., McDirmid, S., Micheloud, S., Mihaylov, N., Schinz, M., Stenman, E., Spoon, L., & Zenger, M. (2006). *An Overview of the Scala Programming Language* (2nd ed.). École Polytechnique Fédérale de Lausanne (EPFL)

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

but are domain experts. Scala’s syntax, type system and abstraction mechanisms make the tool well suited for crafting DSLs. Features — like operator overloading, implicit conversions and higher-order functions — empower developers to craft APIs that feel like natural extensions of the language itself (Odersky et al., 2021)<sup>2</sup>.

Key factors that help a Scala evolve are its **abstraction abilities**. Scala can be used to make patterns and components that can be combined. **Traits**, for instance, allow programmers to define repeatable interfaces and implementations that can be incorporated into classes as needed. **Higher-order functions** and **type classes** help create abstractions that can be used with many types of data. These abstraction mechanisms hold significant importance in the development of libraries and frameworks that can augment the language in novel ways (Odersky et al., 2021)<sup>2</sup>.

Having that said, the idea of **DSLs** wouldn’t make any sense without **community-driven development**. Scala contains a rich ecosystem of libraries and frameworks, of which a few are discussed in the later chapters. These libraries are often made by people who have encountered a problem and decided to use Scala to solve it. Thanks to this design, Scala can stay relevant and adapt to the changing needs of developers or the industry in general (Odersky et al., 2021)<sup>2</sup>.

The design of Scala ensures that **custom libraries feel like native language features**. The unified type system in Scala makes it easy to add new types — like classes, traits or objects — that can be used with other built-in types in pairs. Meaning, libraries can create new types that are as natural as Integers, Strings or Lists (Ghosh, 2011)<sup>3</sup>.

## 1.3 What Makes Scala Scalable?

Scala’s smooth **integration between object-oriented and functional programming** is important for scalability. Object-oriented elements give the structure and flexibility needed for making big and complicated systems. They allow developers to create components that can be combined and expanded as the codebase grows. On the other hand, the functional aspects, support scalability. By combining these two approaches, developers can write code that’s both modular and flexible to meet the changing needs of the system (Odersky et al., 2021)<sup>2</sup>.

In addition, Scala’s ability to use static types is significant for making it **easy to scale**. The **type system** helps find errors early, reduces errors during runtime and makes code more reliable. It supports scalable code evolution through various features that enable the creation of flexible and reusable abstractions. These abstractions can adapt to changing requirements without sacrificing type safety or requiring extensive code rewrites. It is easier to reason about as the codebase grows in size and complexity because the type system encourages developers to express their intent clearly (Odersky et al., 2021)<sup>2</sup>.

Lastly, Scala’s high-level abstractions for data processing and concurrency are the best feature for writing scal-

---

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

able code. **Abstractions** — like **map**, **flatMap** and **fold** (explained later) — allow developers to express complex data transformations and computations in a concise and declarative manner. These abstractions can be used on multiple cores or machines, which makes it easy to process large datasets (Odersky et al., 2021)<sup>2</sup>.

In terms of **concurrency**, Scala has useful tools known as **Futures** and **Promises** that make it easier to create **asynchronous** and **non-blocking code**. Using abstractions and tools like **Akka** help developers to create **responsive** and **resilient systems** that can endure without the need to handle low-level concurrency primitives directly (Odersky et al., 2021)<sup>2</sup>.

## 1.4 Object-Oriented Paradigm in Scala

As previously stated, **Scala is based on the JVM**. Consequently, the code is written in a manner that is similar to that of Java. Object-oriented programming, together with Scala’s functional programming qualifies Scala as a **multi-paradigm language**. In short, as a result, **every value is an object and every operation is a method call** (Suereth, 2012)<sup>4</sup>.

Object-oriented programming offers excellent **modularity** and **encapsulation**, resulting in **reusable code**. **Classes** and **objects** help organize data and behaviors into logical parts. With this approach, code becomes significantly more streamlined and organized. By putting a data into a processing logic in classes and objects, components can be made and used to build **data processing pipelines**. This modular approach makes code easier to maintain and test, which holds great significance in big data projects (Ghosh, 2011)<sup>3</sup>.

In addition, Scala supports **single inheritance through classes**, which allows for hierarchical relationships between data structures and processing logic. Common behaviors and attributes can be defined in base classes and specialized in derived classes, which helps eliminate redundant code and promotes code reuse (Ghosh, 2011)<sup>3</sup>.

**Traits** are a flexible way to combine behavior and define common interfaces. Adding functionality horizontally by mixing it into classes allows for the creation of modular and reusable data processing elements. Common data processing operations — such as **data transformation, aggregation and persistence** — can be defined across different stages of the data pipeline based on inheritance and trait composition (Ghosh, 2011)<sup>3</sup>.

Table 1.1: Scala’s transformation of operators into method calls

```
val x = 1 + 2 * 3 // How it is normally written
val y = 1.+(2.*(3)) // To what the code is transformed under the hood
```

*Note.* This table illustrates Scala’s uniform object model, where arithmetic operators are translated into method invocations on objects. *Creator.* Author’s own work.

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

<sup>4</sup>Suereth, J. D. (2012). *Scala in Depth*. Manning Publications

<sup>3</sup>Ghosh, D. (2011). *DSLs in Action*. Manning Publications

## 1.5 Overview of Functional Programming Paradigm

**Functional programming (FP)** is a way of making programs that uses lambda calculus. Programs are created by combining functions together. In **FP**, computation means evaluating mathematical functions without changing their state or modifying data (Stenberg, 2019)<sup>5</sup>.

Functional programming is based on lambda calculus, a formal system developed by Alonzo Church in the 1930s to study computerability. The first functional programming language, Lisp, was developed by John McCarthy in 1958. Other early influential functional languages include: APL (1966), ML (1973), Scheme (1975), Miranda (1985), Haskell (1990). In the last few years, functional programming has become more popular. Modern multi-paradigm languages like Scala, F#, Clojure and Elixir all support functional programming. Many mainstream languages like Java, C# and Python have also added functional features (Kunasaikaran and Iqbal, 2016)<sup>6</sup>.

## 1.6 Core Features of Functional Programming

The following section provides an overview of **key functional programming concepts** utilized in Scala **for data engineering**. These concepts include immutable data structures, higher-order functions, lazy evaluation, pattern matching, algebraic data types, type classes, monads, error handling and parallel and distributed processing. This thesis does not attempt to provide a comprehensive explanation of functional programming or Scala. Instead, it provides a concise explanation of each concept and offers a straightforward example to demonstrate its usage. Having a solid grasp of these concepts is crucial for building efficient data pipelines. This section focuses on the core principles and how they can be used in data engineering practices.

### 1.6.1 Immutable data structures

**Immutable data structures** are fundamental, in functional programming. It plays a key role, in constructing data pipelines. In Scala, these types of data structures are part of the language and its standard library. An immutable data structure is one that **stays the same after it is created. Once its structure is set, it stays that way all the time**. The original data structure is left unchanged by any action that seems to alter the structure (Zibin et al., 2007)<sup>7</sup>.

As the name "immutable" indicates, **data consistency** makes sure that the data stays the same throughout the data processing pipeline. By not letting changes to the data happen, there is no chance for the data to get corrupted

---

<sup>5</sup>Stenberg, M. (2019). Functional and imperative object-oriented programming in theory and practice. <https://api.semanticscholar.org/CorpusID:197640090>

<sup>6</sup>Kunasaikaran, J., & Iqbal, A. (2016). A brief overview of functional programming languages. *electronic Journal of Computer Science and Information Technology (eJCSIT)*, 6(No. 1). <https://api.semanticscholar.org/CorpusID:195989733>

<sup>7</sup>Zibin, Y., Potanin, A., Ali, M., Artzi, S., Kie'zun, A., & Ernst, M. D. (2007). Object and Reference Immutability using Java Generics. *MIT Computer Science & Artificial Intelligence Lab*



or inconsistent. This is crucial when many processes or threads are working on the data at once (Milewski, 2013)<sup>8</sup>.

Another advantage is that data structures are inherently **thread-safe**, as they cannot be modified by multiple threads simultaneously. This eliminates the necessity for synchronization mechanisms. In addition, it helps to reduce the likelihood of race conditions and other bugs related to concurrency. Data can be easily shared and accessed by multiple threads without the need for locks or other synchronization primitives (Milewski, 2013)<sup>8</sup>.

Another notable benefit is **fault tolerance**. It is the ability to handle errors and continue functioning effectively. Regardless of any circumstances, the data remains consistent. The processing can resume from a recognized point. This feature is incredibly valuable for systems that experience frequent failures and require the ability to recover from setbacks in order to maintain accurate and consistent data (Milewski, 2013)<sup>8</sup>.

Scala has many types of immutable data — like **List**, **Vector**, **Map** and **Set** — in its standard library. These **data structures (DS)** were created to maximize efficiency. Tailored to practical scenarios. For example, List is a **linked-list** implementation that can be used for recursive algorithms and pattern matching. On the other hand, the **Vector DS** is based on trees, which provide impressive random access and updates, making it perfect for indexed data processing (Agarwal, 2023)<sup>9</sup>.

Table 1.2: Immutability and transforming lists

```
val data = List(1, 2, 3, 4, 5)
val updatedData = data.map(_ * 2)
println(data)           // Output: List(1, 2, 3, 4, 5)
println(updatedData)    // Output: List(2, 4, 6, 8, 10)
```

*Note.* In this example, the **data** list is an immutable data structure. The **map** operation creates a new list, **updatedData** by applying the transformation function to each element of **data**. The original **data** list remains unchanged, ensuring data consistency and thread safety. *Creator.* Author's own work.

In addition, when data cannot be changed, it allows for easy sharing between different parts of the data processing pipeline. Since the information stays the same, it can be shared and used again without worrying about copying or synchronizing. It especially shines when dealing with large datasets by improving performance (Tome et al., 2024)<sup>10</sup>.

<sup>8</sup>Milewski, B. (2013). Functional data structures and concurrency in c++. Retrieved June 13, 2024, from <https://bartoszmilewski.com/2013/12/10/functional-data-structures-and-concurrency-in-c/>

<sup>9</sup>Agarwal, Y. (2023). Scala vector. Retrieved June 13, 2024, from <https://www.scaler.com/topics/scala/scalavector/>

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

Table 1.3: Partitioning a vector based on a predicate

```
val data = Vector(1, 2, 3, 4, 5)
val evenData = data.filter(_ % 2 == 0)
val oddData = data.filter(_ % 2 != 0)
println(evenData) // Output: Vector(2, 4)
println(oddData) // Output: Vector(1, 3, 5)
```

*Note.* In this example, the **data** vector is shared between the **evenData** and **oddData** computations. Since **data** is immutable, it can be safely accessed and filtered by both computations without the risk of data races or inconsistencies. *Creator:* Author’s own work.

The fact that data is immutable, results in the ability to create and use **lazy evaluation (LE)** and optimize data processing pipelines. **LE** lets delay the computation of intermediate results until they are needed, which results in reduced memory overhead and improved overall performance. Scala’s data structures — such as **Stream** and **Iterator** — provide **LE** capabilities out of the box, which makes it easy to create memory-friendly data pipelines (Tome et al., 2024)<sup>10</sup>.

Table 1.4: Example of lazy evaluation and infinite sequences

```
val data = Stream.from(1)
val evenData = data.filter(_ % 2 == 0).take(5)
println(evenData.toList) // Output: List(2, 4, 6, 8, 10)
```

*Note.* In this example, the **data** stream is an infinite sequence of integers starting from 1. The **filter** operation lazily filters the even numbers from the stream, and the **take** operation limits the result to the first 5 even numbers. The computation is performed lazily, generating only the required elements on-demand, thus avoiding the need to materialize the entire infinite sequence in memory. *Creator:* Author’s own work.

## 1.6.2 Higher-Order Functions

**Higher-order functions (HOFs)** are yet another killer feature; basically, they help make data processing systems that can be easily changed and reused. Through their utilization, the code can be fine-tuned to be faster and easier to maintain. These functions typically used for data transformation, grouping and analysis (Pilquist et al., 2023)<sup>11</sup>.

The **HOF** can be defined as a type of function that takes one or more functions as inputs, gives back a function as an output, or does both. In more complex terms, this feature of treating **functions as first-class citizens** permits the development of adaptable data processing tasks. These functions help change behavioral parameters and combine functions. (Pilquist et al., 2023)<sup>11</sup>.

In the case of examples, the Scala standard library provides a rich set of tools that are commonly used in data engineering tasks. These functions operate on collections and enable concise and expressive data manipulation.

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

Below are the most popular ones (Pilquist et al., 2023)<sup>11</sup>.

Table 1.5: map

```
val data = List(1, 2, 3, 4, 5)
val squaredData = data.map(x => x * x)
println(squaredData) // Output: List(1, 4, 9, 16, 25)
```

*Note.* The **map** function applies a given function to each element of a collection and returns a new collection with the transformed elements. It is used for element-wise transformations and is a fundamental building block of data processing pipelines. *Creator:* Author's own work.

Table 1.6: flatMap

```
val data = List("hello world", "functional programming", "data engineering")
val words = data.flatMap(_.split(" "))
println(words) // Output: List("hello", "world", "functional", "programming", "data", "engineering")
```

*Note.* The **flatMap** function applies a given function to each element of a collection and flattens the resulting collections into a single collection. It is used for transformations that produce zero or more elements per input element, and is particularly useful for data flattening and joining operations.

Table 1.7: filter

```
val data = List(1, 2, 3, 4, 5)
val evenData = data.filter(_ % 2 == 0)
println(evenData) // Output: List(2, 4)
```

*Note.* The **filter** function selects elements from a collection based on a given predicate function. It is used for data filtering and selection operations, allowing for the creation of subsets of data that satisfy specific criteria. *Creator:* Author's own work.

Table 1.8: reduce

```
val data = List(1, 2, 3, 4, 5)
val sum = data.reduce(_ + _)
println(sum) // Output: 15
```

*Note.* The **reduce** function combines the elements of a collection using a binary operator function. It is used for data aggregation and summarization tasks, such as computing sums, products, or custom aggregations. *Creator:* Author's own work.

As shown in the 1.9 table, **HOFs** enable data processing operations that can be reused regardless of the data type. With the ability to establish the internal structure of systems using functions, one can create data processing pipelines that are adaptable to various data types and requirements (Pilquist et al., 2023)<sup>11</sup>.

**HOFs** also enable the creation of **domain-specific languages (DSLs)** and fluent **APIs** for data processing. By defining a set of **HOFs**, data engineers can make code that looks like the problem domain (Pilquist et al., 2023)<sup>11</sup>.

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

Table 1.9: Generic data processing pipeline with higher-order functions

```
def processData[A, B](data: List[A],
  transformFn: A => B, filterFn: B => Boolean, aggregateFn: (B, B) => B): B = {
  data.map(transformFn)
    .filter(filterFn)
    .reduce(aggregateFn)}
val data = List(1, 2, 3, 4, 5)
val sum = processData(data, x => x * x, x => x > 10, _ + _)
println(sum) // Output: 41
```

*Note.* In this example, the **processData** function is a higher-order function that takes three function parameters: **transformFn** for data transformation, **filterFn** for data filtering, and **aggregateFn** for data aggregation. By providing different functions as arguments, the **processData** function can be easily customized for different data processing scenarios, promoting code reuse and modularity. *Creator:* Author's own work.

Table 1.10: Building a data processing DSL

```
case class DataPipeline[A](data: List[A]) {
  def map[B](f: A => B): DataPipeline[B] = DataPipeline(data.map(f))
  def filter(f: A => Boolean): DataPipeline[A] = DataPipeline(data.filter(f))
  def reduce[B >: A](f: (B, B) => B): B = data.reduce(f)}
val pipeline = DataPipeline(List(1, 2, 3, 4, 5))
  .map(x => x * x)
  .filter(x => x > 10)
  .reduce(_ + _)

println(pipeline) // Output: 41
```

*Note.* In this example, the **DataPipeline** class defines a fluent API for data processing using higher-order functions. The **map**, **filter**, and **reduce** methods enable the creation of expressive and readable data processing pipelines that can be easily composed and chained together. *Creator:* Author's own work.

### 1.6.3 Lazy Evaluation

**Lazy evaluation (LE)** is a technique where the evaluation of an expression is delayed until its value is actually needed — look at table 1.11. It makes big amounts of information work faster, avoids unnecessary work and let make endless data structures (Pawar, 2023)<sup>12</sup>.

In Scala, **LE** is possible by using **Lazy Values (LV)**, **Lazy Collections (LC)** and **lazy data processing**. In terms of **LV**, they are only evaluated when they are accessed for the first time. Later accesses to a **LV** reuse the result that was previously computed (Pilquist et al., 2023)<sup>11</sup>.

<sup>12</sup>Pawar, H. (2023). Scala lazy evaluation. Retrieved June 13, 2024, from <https://www.scaler.com/topics/scala/lazy-evaluation-in-scala/>

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

Table 1.11: Lazy evaluation

```

lazy val data = {
  println("Initializing data...")
  List(1, 2, 3, 4, 5)}
println("Before accessing data")
val result = data.map(_ * 2)
println("After accessing data")

```

*Note.* In this example, the **data** value is defined as a lazy value. The initialization code inside the block is not executed until **data** is accessed for the first time. This allows for the deferred execution of potentially expensive computations and enables the creation of data processing pipelines that can handle large datasets efficiently. *Creator:* Author’s own work.

In the case of **LC**, — such as **Stream** and **Iterator** — a lazy and incremental approach to processing data is provided by these collections, which represent potentially infinite sequences of elements. A **LCs** elements are computed on demand, which makes it easy to process large or infinite datasets (Hughes, 1990)<sup>13</sup>.

Table 1.12: Infinite fibonacci stream

```

def fibonacci: Stream[Int] = 0 #:: 1 #:: fibonacci.zip(fibonacci.tail).map(t => t._1 + t._2)
val fibs = fibonacci.take(10).toList
println(fibs) // Output: List(0, 1, 1, 2, 3, 5, 8, 13, 21, 34)

```

*Note.* In this example, the **fibonacci** function defines an infinite stream of Fibonacci numbers using lazy evaluation. The **#::** operator is used to construct the stream lazily, and the **zip** and **map** operations are used to generate the next Fibonacci number based on the previous two. The **take** operation limits the evaluation to the first 10 elements, avoiding the need to compute the entire infinite sequence. *Creator:* Author’s own work.

Another advantage of **LE** is that it can be employed to construct large data processing systems capable of handling large amounts of information, such as terabytes or petabytes. Furthermore, **LE** can be employed to process data incrementally, focusing solely on the essential computations (Chen et al., 2011)<sup>14</sup>.

In the case of Scala’s functional programming libraries, — such as **Apache Spark** and **Akka Stream**; table 1.12 — they use **LE** a lot to process data in many places at once. In **Apache Spark**, **LE** is used to build a **directed acyclic graph (DAG)** — as can be seen in the 1.13 table — of transformations and actions that are optimized and executed in a distributed way across a cluster of machines. **Spark** can use **LE** to make the execution plan work better, move data less often and improve performance overall (Pilquist et al., 2023)<sup>11</sup>.

<sup>13</sup>Hughes, J. (1990). Why Functional Programming Matters. “*Research Topics in Functional Programming*” ed. D. Turner, Addison-Wesley, 17–42

<sup>14</sup>Chen, G., Chen, K., Jiang, D., Ooi, B. C., Shi, L., Vo, H. T., & Wu, S. (2011). E3 : An Elastic Execution Engine for Scalable Data Processing. *Journal of Information Processing*, 20(No.1), 65–76

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

Table 1.13: Direct acrylic graph (DAG)

```
val data = spark.textFile("hdfs://path/to/data.txt")
val words = data.flatMap(_.split(" "))
val wordCounts = words.map(_._1).reduceByKey(_ + _)
wordCounts.saveAsTextFile("hdfs://path/to/output")
```

*Note.* In this example, the **data** RDD (Resilient Distributed Dataset) represents a large text file stored in HDFS. The **flatMap**, **map**, and **reduceByKey** operations are lazily evaluated, building a DAG of transformations. The actual computation is triggered only when the **saveAsTextFile** action is called, allowing Spark to optimize the execution plan and distribute the processing across the cluster. *Creator:* Author's own work.

By utilizing lazily defined data transformations and computations, data engineers can construct reusable and adaptable components that can be easily combined and adapted to diverse data processing scenarios — look at table 1.14. This strategy encourages code reuse, improves upkeep and aids in the development of intricate data processing workflows (Pilquist et al., 2023)<sup>14</sup>.

Table 1.14: Data processing workflow

```
def loadData(path: String): Stream[String] = {
  val source = io.Source.fromFile(path)
  val lines = source.getLines().toStream
  source.close()
  lines}

def processData(data: Stream[String]): Stream[(String, Int)] = {
  data.flatMap(_.split(" "))
    .map(_._1)
    .groupBy(_._1)
    .mapValues(_._2.sum)
    .toStream
}

val data = loadData("path/to/data.txt")
val processed = processData(data)

processed.take(10).foreach(println)
```

*Note.* In this example, the **loadData** function lazily loads data from a file and returns a stream of lines. The **processData** function defines a series of lazy transformations on the input stream, including splitting lines into words, counting word occurrences, and grouping by word. The **take** operation triggers the evaluation of the first 10 elements, and the **foreach** operation prints the results. This modular and composable approach allows for the creation of reusable data processing components that can be easily combined and adapted to different datasets and requirements. *Creator:* Author's own work.

<sup>14</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

## 1.6.4 Pattern Matching and Algebraic Data Types (ADTs)

**Pattern matching** and **algebraic data types** — for shorts, **PM** and **ADTs** — provide an expressive way to model, process and retrieve information from complex data structures. **PM** makes code clear and easy to understand and **ADTs** make data structures that are expressive and type-safe — look at table 1.15 (Pilquist et al., 2023)<sup>11</sup>.

Despite the complex name, **ADTs** are just a way to define complex data structures using a combination of product types — such as **tuples** or **case classes** — and summarize types — like **sealed traits** or **enums** —. Product types are data with multiple fields, while summarize types are data with only one field. What makes **ADTs** useful is that they allow for the modeling of hierarchical and recursive data structures (Pilquist et al., 2023)<sup>11</sup>.

They are typically defined using **sealed traits** and **case classes**. The first ones define a set of possible variants that covers all cases and prevents the addition of new variants outside the set. Case classes, on the other hand, represent the individual variants of the **ADT** and provide a convenient way to define and manipulate data (Pilquist et al., 2023)<sup>11</sup>.

Table 1.15: Data model using ADTs

```
sealed trait Shape
case class Circle(radius: Double) extends Shape
case class Rectangle(width: Double, height: Double) extends Shape
case class Triangle(base: Double, height: Double) extends Shape
```

*Note.* In this example, the **Shape** ADT represents different geometric shapes. It is defined as a sealed trait, and the individual shapes (Circle, Rectangle, Triangle) are defined as case classes that extend the **Shape** trait. This allows for the creation of a type-safe and expressive data model for representing shapes. *Creator:* Author's own work.

As mentioned before, the **PM** is a feature in Scala that allows for the concise and expressive processing of **ADTs** — look at table 1.16 —. Basically, it enables the deconstruction and extraction of data based on its structure and variant (Pilquist et al., 2023)<sup>11</sup>.

Table 1.16: Pattern matching

```
def area(shape: Shape): Double = shape match {
  case Circle(radius) => math.Pi * radius * radius
  case Rectangle(width, height) => width * height
  case Triangle(base, height) => 0.5 * base * height}
```

*Note.* In this example, the **area** function uses pattern matching to calculate the area of different shapes. The **match** expression checks the type of the **shape** parameter and extracts the relevant fields based on the corresponding case class. This allows for concise and expressive code that handles each shape variant separately. *Creator:* Author's own work.

**PM** can also be used to extract and reorganize data. It is done by using patterns and conditions to get data that needs to be filtered, transformed or grouped based on certain criteria — look at table 1.17 (Pilquist et al., 2023)<sup>11</sup>.

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications

Table 1.17: Data processing and aggregation

```
def processData(data: List[Shape]): Double = data match {
  case Nil => 0.0
  case Circle(radius) :: rest => radius * radius + processData(rest)
  case Rectangle(width, height) :: rest => width * height + processData(rest)
  case Triangle(base, height) :: rest => 0.5 * base * height + processData(rest)}
```

*Note.* In this example, the **processData** function uses pattern matching to process a list of shapes. It recursively matches on the head of the list and extracts the relevant fields based on the shape variant. The function calculates a value based on the shape and recursively processes the rest of the list. This demonstrates how pattern matching can be used for data processing and aggregation. *Creator:* Author's own work.

For example, when a data engineer has to work with complex and hierarchical data structures. He can create type-safe data models that can be easily processed and transformed — look at table 1.18 (Pilquist et al., 2023)<sup>11</sup>.

Table 1.18: Pattern matching and ADTs

```
case class User(id: Int, name: String, age: Int)
case class Transaction(userId: Int, amount: Double, timestamp: Long)

def processTransactions(transactions: List[Transaction]): Map[Int, Double] = transactions match {
  case Nil => Map.empty[Int, Double]
  case Transaction(userId, amount, _) :: rest =>
    val userTotal = processTransactions(rest).getOrElse(userId, 0.0) + amount
    processTransactions(rest) + (userId -> userTotal)
}

val users = List(User(1, "Alice", 25), User(2, "Bob", 30), User(3, "Charlie", 35))

val transactions = List(
  Transaction(1, 100.0, System.currentTimeMillis()),
  Transaction(2, 50.0, System.currentTimeMillis()),
  Transaction(1, 75.0, System.currentTimeMillis()),
  Transaction(3, 120.0, System.currentTimeMillis())

val userTotals = processTransactions(transactions)
val userSummary = users.map(user => (user, userTotals.getOrElse(user.id, 0.0)))
```

*Note.* In this example, ADTs is defined for **User** and **Transaction**. The **processTransactions** function uses pattern matching to process a list of transactions and calculate the total amount for each user. The **userSummary** combines the user information with their transaction totals using pattern matching and the **map** operation. This demonstrates how pattern matching and ADTs can be used together to process and analyze complex data structures in a data engineering context. *Creator:* Author's own work.

<sup>11</sup>Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications



## 1.6.5 Type Classes

In functional programming, **type classes** — for short, **TC** — are an idea that allows for ad-hoc polymorphism and offers a way to define generic behavior for types without altering their primary definitions. These classes are crucial as they facilitate the development of reusable and composable abstractions that're applicable, to various data types and structures (Odersky et al., 2021)<sup>2</sup>. **TCs** are implemented using implicit parameters and implicit definitions — look at table 1.19 —; a **TC** is defined as a trait that declares a set of operations or behaviors that can be implemented for different types. These operations are defined as abstract methods within the trait. Types that want to belong to the **TCs** provide implicit implementations of these methods (Odersky et al., 2021)<sup>2</sup>.

Table 1.19: Implicit implementations

```
trait Semigroup[A] {  
  def combine(x: A, y: A): A  
object Semigroup {  
  implicit val intSemigroup: Semigroup[Int] = new Semigroup[Int] {  
    def combine(x: Int, y: Int): Int = x + y  
  }  
  implicit val stringSemigroup: Semigroup[String] = new Semigroup[String] {  
    def combine(x: String, y: String): String = x + y  
  }  
}
```

*Note.* In this example, a **Semigroup** type class is defined that represents types with an associative binary operation **combine**. Implicit implementations of **Semigroup** for **Int** and **String** types are provided, defining the **combine** operation as addition and concatenation, respectively. *Creator.* Author's own work.

**TCs** help create generic functions — look at table 1.20 —; These are the functions that can be used with any type that has an implicit instance of the **TC** needed. In other words, it makes it possible to make abstractions that can be used on different data types without changing their original definitions (Odersky et al., 2021)<sup>2</sup>.

Table 1.20: Generic function

```
def combineAll[A](values: List[A])(implicit semigroup: Semigroup[A]): A = values.reduce(semigroup.combine)  
val numbers = List(1, 2, 3, 4, 5)  
val strings = List("Hello", " ", " ", "world", "!")  
println(combineAll(numbers)) // Output: 15  
println(combineAll(strings)) // Output: "Hello, world!"
```

*Note.* In this example, the **combineAll** function takes a list of values of type **A** and an implicit **Semigroup** instance for type **A**. It uses the **combine** operation provided by the **Semigroup** to reduce the list of values into a single value. The function can be called with any type that has an implicit **Semigroup** instance, such as **Int** and **String**, without requiring any modifications to the function itself. *Creator.* Author's own work.

**TCs** are handy when there is a need to define common behaviors or operations for different data types — despite lots of code table 1.21 explains it pretty well —; They can be used to make data processing abstractions that can be used in a wide range of data structures (Odersky et al., 2021)<sup>2</sup>.

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

Table 1.21: Implicit instance

```

trait Encoder[A] {def encode(value: A): String}
object Encoder {
  implicit val intEncoder: Encoder[Int] = new Encoder[Int] {
    def encode(value: Int): String = value.toString
  }
  implicit val stringEncoder: Encoder[String] = new Encoder[String] {
    def encode(value: String): String = value
  }
  implicit def listEncoder[A](implicit encoder: Encoder[A]): Encoder[List[A]] = new Encoder[List[A]] {
    def encode(values: List[A]): String = values.map(encoder.encode).mkString("[", ", ", "]")
  }
  def processData[A](data: List[A])(implicit encoder: Encoder[A]): String =
    data.map(encoder.encode).mkString(",")
  val numbers = List(1, 2, 3, 4, 5)
  val strings = List("apple", "banana", "orange")
  val nested = List(List(1, 2), List(3, 4), List(5, 6))
  println(processData(numbers)) // Output: "1,2,3,4,5"
  println(processData(strings)) // Output: "apple,banana,orange"
  println(processData(nested)) // Output: "[1,2],[3,4],[5,6]"
}

```

*Note.* In this example, an **Encoder** type class is defined, it provides an **encode** operation to convert values of type **A** to a string representation. Implicit instances of **Encoder** for **Int** and **String** types are provided. Also, an implicit **listEncoder** is defined, it can encode a list of values of type **A**, given an implicit **Encoder** instance for type **A**. The **processData** function takes a list of values of type **A** and an implicit **Encoder** instance for type **A**. It uses the **encode** operation to convert each value to its string representation and concatenates them into a single string. The function can be called with any type that has an implicit **Encoder** instance, including nested lists, demonstrating the composability and reusability of type classes. *Creator.* Author's own work.

In terms of Scala libraries and **TCs** the Cats and Scalaz are the **TC** are also commonly used in Scala's functional programming libraries, such as Cats and Scalaz, which provide a wide range of type classes for various algebraic structures and data processing operations. These libraries leverage type classes to define generic and composable abstractions for data processing — such as Functor, Monad and Traverse — which can be applied to different data types and structures (Odersky et al., 2021)<sup>2</sup>.

## 1.6.6 Monads and error handling

One of the hardest to understand common concept in Functional programming are **Monads (M)**. In essence they can be used to make and chain data processing operations, handle errors or validate data. Basically an **M** is a design pattern that allows for the composition of computations in a way that is both expressive and type-safe. In simpler terms, it is a structure that wraps a value (or computation) and provides a standardized way to chain operations on that value (Wadler, 1992)<sup>15</sup>.

The two fundamental operations are **flatMap** — also known as **bind** — and **pure** — also known as **unit** or **return** —. These operations follow certain rules called the monadic laws; basically, they make sure that they

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

<sup>15</sup>Wadler, P. (1992). Monads for functional programming

behave the same way (Wadler, 1992)<sup>15</sup>.

In Scala, **M**s are implemented using the **flatMap**, **map** and other **pure methods**. A function can be called **pure when it produces the same output for a given input, no matter what**. In this context, the **flatMap** binds computations that may produce an **M** instance, while the **map** method applies a function to the value inside the **M**. A value is inserted into the **M** context by the pure approach (Suereth, 2012)<sup>4</sup>.

In the Scala's standard library there are **M** implementations for the most common types — such as **Option**, **Either**, **Try** and **Future**, each serving different purposes in error handling and data processing. The most popular scenario for using **M**s is data validation and error handling. **M**s — like **Option** and **Either** — can be used to represent errors or missing values in a data pipeline. **Option** is used to represent a value that may or may not be present, while **Either** represents a value that can be either a success or a failure. In result the data pipeline remains in a consistent state and avoids unexpected failures (Suereth, 2012)<sup>4</sup>.

For example, lets think about a data pipeline that involves typical ETL scenario. Each step in the process can go wrong — like there is a missing file, record or a bug in transformation —; To solve these problems, **Either** can be used for centralized error handling and reporting by capturing these errors through the pipeline (Suereth, 2012)<sup>4</sup>.

Table 1.22: Monads and error handling

```
import scala.io.Source

def readFile(filename: String): Either[String, List[String]] =
  try {val lines = Source.fromFile(filename).getLines().toList
      Right(lines)} catch {case e: Exception => Left(s"Error reading file: ${e.getMessage}")}

def parseData(lines: List[String]): Either[String, List[Int]] =
  try {val parsed = lines.map(_.toInt)
      Right(parsed)} catch {case e: NumberFormatException => Left(s"Error parsing data: ${e.getMessage}")}

def processData(data: List[Int]): Either[String, Int] =
  try {val result = // Perform data processing
      Right(result)} catch {case e: Exception => Left(s"Error processing data: ${e.getMessage}")}

val pipeline = for {
  lines <- readFile("data.txt")
  parsed <- parseData(lines)
  result <- processData(parsed)
} yield result

pipeline match {
  case Right(result) => println(s"Data processing successful: $result")
  case Left(error)   => println(s"Data processing failed: $error")}
```

*Note.* In this example, the **readFile**, **parseData**, and **processData** functions return an **Either** monad, representing either a successful result or an error message. The **pipeline** is constructed using a for-comprehension, which chains the operations together using the **flatMap** and **map** methods of **Either**. If any step in the pipeline fails, the error is propagated, and the final result is a **Left** containing the error message. If all steps succeed, the final result is a **Right** containing the processed data.

*Creator:* Author's own work.

<sup>15</sup>Wadler, P. (1992). Monads for functional programming

<sup>4</sup>Suereth, J. D. (2012). *Scala in Depth*. Manning Publications

## 1.6.7 Parallel and Distributed Processing

**Parallel and distributed processing** — **PDP** — in essence, responsible for the efficient processing of large volumes of data. Scala with its libraries and frameworks, is the killer language for **PDP**. Immutable data structures make it safe to share and process data without the risk of race conditions or inconsistencies. Pure functions are inherently parallelizable and can be safely executed independently on multiple nodes in a distributed system (Tome et al., 2024)<sup>10</sup>.

The Parallel Collections library provides a base for parallel processing in Scala. It enables the parallel execution of operations on collections. They can be used to parallelize data processing tasks — such as mapping, filtering and aggregation, table 1.23 — across multiple CPU cores. In result higher performance, especially on big data is achieved. Parallel processing can make a core work faster, but it's better to use more than one core to process a high volume of data (Tome et al., 2024)<sup>10</sup>.

Table 1.23: Parallel vector

```
import scala.collection.parallel.immutable.ParVector
val data: ParVector[Int] = ParVector.range(1, 1000000)

val processed = data.map(_ * 2).filter(_ % 3 == 0).reduce(_ + _)
```

*Note.* In this example, the data is a parallel vector containing a large number of integers. The **map**, **filter**, and **reduce** operations are executed in parallel, utilizing multiple CPU cores to process the data efficiently. *Creator:* Author's own work.

Scala has a rich ecosystem of libraries and frameworks that facilitate distributed data processing, such as Apache Spark and Akka. The first one provides high-level APIs for processing large datasets across clusters of machines. Spark's **Resilient Distributed Datasets** — **RDD** — and **DataFrames** — **DF** — allow for the distributed storage and processing of structured and unstructured data — table 1.24 (Tome et al., 2024)<sup>10</sup>.

Table 1.24: Apache Spark data processing

```
import org.apache.spark.sql.SparkSession
val spark = SparkSession.builder().appName("DataProcessing").getOrCreate()
val data = spark.read.textFile("hdfs://path/to/data.txt")
val processed = data.flatMap(_.split(" ")).map((_, 1)).reduceByKey(_ + _).collect()
```

*Note.* In this example, **Spark** is used to process a large text file stored in HDFS. The **flatMap** operation splits each line into words, the **map** operation converts each word into a tuple of (**word**, **1**), and the **reduceByKey** operation counts the occurrences of each word. The **collect** operation brings the results back to the driver program. *Creator:* Author's own work.

The second framework is **Akka**. It has an actor model that provides a high-level abstraction for **concurrent and distributed computing**, which means that a **scalable** and **fault-tolerant** system can be made of them. Basically, actors are lightweight, independent entities that communicate through **asynchronous message passing**, enabling the distribution of data processing tasks across a cluster of nodes (Tome et al., 2024)<sup>10</sup>.

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

Table 1.25: Akka actors

```
import akka.actor.{Actor, ActorSystem, Props}
case class ProcessData(data: List[Int])
case class ProcessedData(result: Int)
class DataProcessor extends Actor {
  def receive = {case ProcessData(data) => val result = data.map(_ * 2).filter(_ % 3 == 0).sum
    sender() ! ProcessedData(result)}}
val system = ActorSystem("DataProcessingSystem")
val processor = system.actorOf(Props[DataProcessor], "processor")
val data = List(1, 2, 3, 4, 5)
processor ! ProcessData(data)
```

*Note.* In this example, an Akka actor system is created, and a **DataProcessor** actor is defined. The actor receives a **ProcessData** message containing a list of integers, processes the data, and sends back a **ProcessedData** message with the result. The actor can be distributed across multiple nodes in a cluster, allowing for the parallel and distributed processing of data. *Creator.* Author's own work.

## Chapter 2

# Data Engineering Workflow

In the second chapter, a detailed overview of the workflow for data engineering is presented. This overview covers important areas of data management, processing, storage and governance. Beginning with an examination of the dimensions of data quality — its correctness, completeness, consistency, timeliness and uniqueness — the chapter moves on to discuss how to guarantee reliable and valuable data assets.

Following that, the chapter explores data ingestion, processing and storage processes. In it, several data pipeline designs are discussed and a comparison is made between ETL (Extract, Transform, Load) and ELT (Extract, Load, Transform) techniques. Additionally, the benefits and implementation scenarios of each strategy are highlighted.

A considerable section of the chapter is devoted to the discussion of data storage systems, which includes an explanation of the functions that data warehouses, data lakes and NoSQL databases play in contemporary data management. In managing big amounts of different data types and the need for scalability, performance and flexibility.

Additionally covered in this chapter are important facets of data security and governance, including access control techniques, encryption techniques and the difficulties presented by changing legal requirements. It discusses cutting edge data management technologies and methods, such as the use of ethical AI techniques and machine learning for data quality management.

Data integration, serving and the value of monitoring and maintenance in guaranteeing the seamless functioning of data pipelines are covered in the last section. Throughout, the book offers practical examples and addresses the difficulties and possible fixes in putting into practice efficient data engineering techniques.

With its emphasis on its vital role in allowing data-driven decision-making and insights across many sectors, this thorough introduction provides a basis for comprehending the complicated terrain of contemporary data engineering.

## 2.1 Data Quality Dimensions

**Data quality dimensions** — quality and usability of data within a data pipeline.

Before diving into data processing frameworks, it should be clarified on the basis of what data quality is assessed. This subchapter focuses only on 5 data quality dimensions that are the most related to the thesis. However, in the cited paper, there are teens of dimension. Also, it could have been a lot to say about data validation, profiling, cleansing, enrichment, transformations and frameworks, but this is beyond the scope of this master thesis (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

**Accuracy** — how well data represents reality. It measures the proximity of how close the values in the data are in comparison to the prediction. Inaccurate data may cause incorrect analyses, flawed decision-making and unreliable insights. An excellent example would be a customer database. The recorded addresses and contact information have to be accurate and up to date (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

**Completeness** — the amount of information available for analysis. It makes sure that the dataset is not corrupted and/or doesn't have any missing values, fields or records. If the information is not complete, it can make it hard to analyze and give wrong or biased results. For example, in a sales transaction dataset, making sure that all fields, such as customer ID, product ID, quantity and price, are filled in for each transaction (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

**Consistency** — keeping data consistent across different sources, systems and time frames. This makes sure that information is presented in a standard way and follows certain rules for the business. Inconsistent data can lead to confusion, errors and problems when integrating and analyzing data. One way to avoid confusion and make data easier to understand is to use a standard date format called ISO8601 "YYYY-MM-DD" across all data sources (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

**Timeliness** — gauges how current and easily available data is when it comes time for analysis or decision-making. It guarantees current data that accurately depicts the most recent status of the people or events it represents. Outdated or stale data might result in wrong conclusions and suboptimal decisions. One such system may be a real-time fraud detection system, which would guarantee that transaction data is processed and analyzed in near-real time to spot and stop fraudulent activity early on (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

**Uniqueness** — there are no duplicates in a set of data. It guarantees that every record or entity is distinct and consistent. Duplicate information can skew analysis results, consume storage space and lead to incorrect aggregations. In a CRM system, one can use deduplication techniques to find and combine duplicate customer records using unique identifiers or matching criteria (Andrew Black and Nederpelt, 2020)<sup>19</sup>.

---

<sup>19</sup>Andrew Black & Nederpelt, P. van. (2020). Dimensions of Data Quality (DDQ). *DAMA NL Foundation*

## 2.2 Data Ingestion

According to (Sree Sandhya Kona, 2023)<sup>16</sup> **data ingestion**, for short **DI** can be defined as process of collecting and importing data from various sources into a data storage system or processing pipeline. This means getting information from different places — like databases, files, streaming platforms and IoT devices — and using them to analyze it (Sree Sandhya Kona, 2023)<sup>16</sup>.

**DI** frequently handles large volumes of data generated at high speeds, such as real-time streaming data or high-throughput transactional systems. Processing and ingesting such data necessitate scalable **DI** frameworks. The data can appear in multiple formats: structured (e.g., CSV, JSON), semi-structured (e.g., XML, log files) or unstructured (e.g., text, images, videos). **DI** systems must tackle these varying formats effectively to extract useful information. Moreover, these data pipelines must be failure-resilient and uphold data reliability, maintaining stability against network disruptions, system failures and data inconsistencies without losing essential data or compromising its integrity (Sree Sandhya Kona, 2023)<sup>16</sup>.

**Batch Ingestion** involves gathering and importing data at regular intervals into the system. Apache Spark is an example of a distributed processing framework that facilitates batch ingestion, handling large datasets in parallel across a cluster of machines. **Real time streaming** deals with capturing and processing data as it arrives, enabling near-real time analysis and decision-making. These pipelines can be built with tools like **Apache Kafka**, **Apache Flink** or **Akka Streams** (Sree Sandhya Kona, 2023)<sup>16</sup>.

## 2.3 Data Processing

Getting raw data from a lot of different sources, both inside and outside the company, is the first step in **data processing**. From transactional databases and operational systems to online logs, social media feeds and Internet of Things sensor networks, these sources may make up a wide variety of information. Due to the fact that the data may be structured, semi-structured or unstructured, it is necessary to have a robust and scalable ingestion framework that is able to manage a variety of data formats and volumes. Data engineers need to come up with and set up efficient ways to get this data out of storage and into a central repository or data lake so it can be processed further (Tome et al., 2024)<sup>10</sup>.

There are a lot of problems with **raw data**, — like missing values, duplicates, inconsistencies and errors — which can have a big effect on the quality and reliability of analysis that comes after. The goal of the data cleaning phase is to find and fix these problems. It does this by removing duplicate data, dealing with missing values by adding them or deleting them, changing the data type and finding and fixing outliers. This process makes sure that

---

<sup>16</sup>Sree Sandhya Kona. (2023). Leveraging spark and pyspark for data-driven success: Insights and best practices including parallel processing, data partitioning, and fault tolerance mechanisms. *Journal of Mathematical & Computer Applications*. <https://api.semanticscholar.org/CorpusID:270301015>

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing



the data meets quality standards that have already been set and is ready to be transformed and analyzed (Tome et al., 2024)<sup>10</sup>.

Following the completion of the data cleaning and validation processes, it is then subjected to a transformation process in order to be converted into a format that is optimized for analysis and storage. Data normalization, aggregation, joining data from different sources, feature engineering and data enrichment are some of the tasks that may be done at this stage. The changed data is usually organized in a denormalized or dimensional format, which makes it easier to query and analyze. When designing and putting these transformation pipelines into action, data engineers need to be very careful to think about things like performance, scalability and maintainability (Tome et al., 2024)<sup>10</sup>.

Scientists and data analysts can use the processed and changed data in a variety of ways to find insights and help people make decisions based on data. This could include things like data mining, machine learning, statistical analysis and data visualization. These people are closely involved with data engineers, who make sure that the processed data meets their needs and can be easily added to their analytical workflows (Tome et al., 2024)<sup>10</sup>.

After the data has been processed and converted, the last step in the data processing process is to store the data in a format that is optimized for querying and for usage in the future. Depending on the needs of the business, this usually means putting the data into data warehouses, data lakes or NoSQL systems. Data engineers are responsible for designing and implementing storage schemas, indexing algorithms and partitioning approaches in order to guarantee optimum performance, scalability and cost-efficiency for analytical workloads they are responsible for (Tome et al., 2024)<sup>10</sup>.

## 2.4 Data Storage Systems

People in the business world today think of data as the energy that keeps companies going. Companies nowadays are aware that data plays an essential part in their business intelligence systems, which enables them to make choices based on accurate information and to keep a competitive edge in their particular industries (Nambiar and Mundra, 2022)<sup>20</sup>. The term "big data" is widely used to refer to the huge amounts of data that are produced, collected and processed by businesses. Managing and studying this data well is necessary to get useful insights and help people make decisions based on data (Nambiar and Mundra, 2022)<sup>20</sup>. Traditional ways of storing data aren't able to keep up with the huge amounts and types of data that are being created. It can be hard and complicated to deal with and look at the huge amount and variety of "big data." A company can find useful information that helps shape its business plans, though, if it knows how to use and benefit from its huge data collection (Nambiar and Mundra, 2022)<sup>20</sup>. To deal with these problems, businesses need data storing options that are both efficient and flexible. For the purpose of storing and managing large amounts of data, data warehouses and data lakes have become more

---

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

<sup>20</sup>Nambiar, A. M., & Mundra, D. (2022). An overview of data warehouse and data lake in modern enterprise data management. *Big Data Cogn. Comput.*, 6, 132. <https://api.semanticscholar.org/CorpusID:253428554>

popular as data management systems (Nambiar and Mundra, 2022)<sup>20</sup>.

Both **data warehouses (DW)** and **data lakes (DL)** are centralized repositories that are used for the purpose of storing and managing the data used by a business (Nambiar and Mundra, 2022)<sup>20</sup>. Although there are some parallels between **DW** and **DL**, there are also differences between them in terms of their properties and applications. In addition to this, they possess unique qualities that make them suited for a variety of data storage and analysis applications (Nambiar and Mundra, 2022)<sup>20</sup>. **DW** saves data in an ordered and organized way, using predefined schemas. They are made for organized data and are built for quick analysis and queries. **DL**, on the other hand, saves data in its original state, which means that structured, semi-structured and unstructured data can all be stored (Nambiar and Mundra, 2022)<sup>20</sup>. Another thing is how efficient it is. A relational model has problems with scaling, which means that as the amount of data grows, speed drops quickly. As an answer, a new data model was used. **NoSQL** databases, like document stores, graph databases, key-value stores and column-oriented databases, were made as replacements to relational databases. They make it easier to handle big data by giving you more scalability, flexibility and speed (Nayak, 2013)<sup>21</sup>.

The basis for storing, managing and accessing data at different stages of the pipeline is provided by data storage systems, which play a significant role in the pipeline-based data management system. The **scalability factor** is very important for showing how important these systems are. **Data storage** systems need to be able to scale horizontally to meet the growing storage needs as data amounts continue to grow at an exponential rate. Distributed storage options like Amazon S3, Apache Hadoop HDFS and others are made to scale by adding more nodes to the cluster. This makes it possible to handle very large datasets quickly. Even in the event that the hardware fails, these systems include fault tolerance and data replication procedures to guarantee the availability and longevity of the data contained inside them. Scalable data storage systems also allow the creation of data pipelines that can meet the ever-growing needs of big data applications. This makes it easier to store and analyze **terabytes or petabytes of data**. It is very important to be able to flexibly **scale storage space** without any downtime or performance problems in order to keep performance and reliability high while dealing with growing data amounts and complexity in data flows (Nambiar and Mundra, 2022)<sup>20</sup>.

## 2.4.1 The Role and Challenges Associated with Storing

It is during **ingestion phase** that data storage systems receive data from a variety of sources, such as databases, files, APIs and streaming services. They have to keep track of the huge amount of data and how quickly it moves so that storing is reliable and effective. At this stage, it's hard to deal with scalability because storage systems need to **grow to hold more data without losing speed or data integrity**. Distributed storage systems like Amazon S3 and Apache Hadoop HDFS are made to grow horizontally by adding more nodes to the cluster. This lets big data to be stored and processed (Nambiar and Mundra, 2022)<sup>20</sup>.

---

<sup>21</sup>Nayak, A. (2013). Type of nosql databases and its comparison with relational databases. <https://api.semanticscholar.org/CorpusID:15594476>

<sup>20</sup>Nambiar, A. M., & Mundra, D. (2022). An overview of data warehouse and data lake in modern enterprise data management. *Big Data Cogn. Comput.*, 6, 132. <https://api.semanticscholar.org/CorpusID:253428554>

During the **processing stage**, data storage systems are both the source and the destination for transformations and calculations that are made on data. They need to make it easy to get to saved data quickly so that data processing tools can get information and change it effectively. At this point, performance is vital because **storage systems need to be able to handle large amounts of data quickly and with low latency**. By spreading data and calculations across many nodes in the cluster, distributed storage systems and parallel processing tools like Apache Spark make it possible to work with big datasets (Nambiar and Mundra, 2022)<sup>20</sup>.

Another big problem with storing and handling a large scale data is making sure that **the data is consistent**. When many systems and users are viewing and changing the data at the same time, it is paramount to make sure that it is consistent and correct. To keep data consistent, data store systems need to have ways to do it. For example, **ACID (Atomicity, Consistency, Isolation and Durability) traits** in standard databases or eventual **consistency models** in distributed systems are two examples. Data splitting, replication and versioning are all techniques that help keep data consistent and allow for fault tolerance in case of breakdowns or network partitions (Nambiar and Mundra, 2022)<sup>20</sup>.

During the **data serving phase**, data storage systems are what send processed and aggregated data to users and apps further downstream. They need to be able to query and retrieve data quickly and easily, often in real time or near real time. At this stage, **scalability and speed are vital** due to the fact data storing systems need to be able to handle a lot of queries at the same time and return answers with low latency. Distributed databases, including **Apache Cassandra** and **Apache HBase**, are intended to offer high-performance data serving capabilities that are both scalable and performant, thereby facilitating the rapid access to large datasets (Nambiar and Mundra, 2022)<sup>20</sup>.

## 2.4.2 Types of Data Storage Systems

**Data warehouses** are centralized repositories that are used for the purpose of storing and managing massive amounts of structured data that originate from a variety of sources. As a result of their ability to enable complicated queries, data processing and reporting, they make it possible for enterprises to acquire useful insights from their data (Nambiar and Mundra, 2022)<sup>20</sup>.

1. A **schema-on-write method** is what they usually use. This means that the data structure is set up front and data is changed and checked before it is put into the warehouse. This makes sure that the data is correct and consistent because it follows a set format (Nambiar and Mundra, 2022)<sup>20</sup>;
2. Data warehouses are designed to handle **Online Analytical Processing (OLAP)** tasks like complicated searches, data consolidation and multiple analysis. They have features like sorting, splitting and materialized views that make it easier to query and analyze big datasets quickly (Nambiar and Mundra, 2022)<sup>20</sup>;
3. They are made to **handle complicated queries** that involve joining multiple tables, collecting data and doing math. They have query tools that are built on SQL and support advanced query optimization methods to make

---

<sup>20</sup>Nambiar, A. M., & Mundra, D. (2022). An overview of data warehouse and data lake in modern enterprise data management. *Big Data Cogn. Comput.*, 6, 132. <https://api.semanticscholar.org/CorpusID:253428554>

sure that queries run quickly. Amazon Redshift, Google BigQuery and Snowflake are all types of data centers (Nambiar and Mundra, 2022)<sup>20</sup>.

Another part of data storage systems are **data lakes**. They let companies use the power of big data analytics by giving them a central place to store huge amounts of raw, unstructured and semi-structured data from different sources (Nambiar and Mundra, 2022)<sup>20</sup>.

1. The way that data lakes and data warehouses handle schemas is different. A schema-on-write method is used in data warehouses, while a **schema-on-read** method is used in data lakes. In other words, data that is put into a data lake is kept in its original form, without a model that has already been set up. The schema is used when the data is viewed and handled, so it can be changed as the needs of the data change (Nambiar and Mundra, 2022)<sup>20</sup>;
2. **Different types of data** can be stored in data lakes. This includes structured data (like CSV and JSON), semi-structured data (like XML and log files) and unstructured (like pictures and videos). Because of this, businesses can get information from a variety of places, such as social media sites, Internet of Things (IoT) devices and clickstream datasets (Nambiar and Mundra, 2022)<sup>20</sup>;
3. Data lakes are a way to **store large amounts of data for a low cost**. In order to achieve scalability, robustness and cost effectiveness, they make use of cloud services such as Amazon S3 and Azure Data Lake Storage, as well as distributed storage solutions such as Apache Hadoop HDFS. The data can be stored on hardware or in the cloud and these systems make sure that it is always available. Apache Hadoop HDFS, Amazon S3 and Azure Data Lake Storage are all well-known examples of data lakes (Nambiar and Mundra, 2022)<sup>20</sup>.

**NoSQL databases** effective handling of semi-structured and unstructured data has made them quite popular in recent years. Because they provide a scalable and versatile substitute for conventional relational databases, they are ideal for managing the wide range of quickly expanding data needs of contemporary data pipelines (Nayak, 2013)<sup>21</sup>.

1. **Schema-less** or **schema-flexible** databases enable data to be stored without a specified schema. This flexibility makes it possible to quickly design and iterate data-driven applications and to easily accommodate changing data structures (Nayak, 2013)<sup>21</sup>;
2. Because they are **designed to grow horizontally**, NoSQL databases may have data spread across many cluster nodes. Because they provide low latency and excellent speed for read and write operations, they are appropriate for managing big data volumes and high-throughput applications (Nayak, 2013)<sup>21</sup>;

---

<sup>20</sup>Nambiar, A. M., & Mundra, D. (2022). An overview of data warehouse and data lake in modern enterprise data management. *Big Data Cogn. Comput.*, 6, 132. <https://api.semanticscholar.org/CorpusID:253428554>

<sup>21</sup>Nayak, A. (2013). Type of nosql databases and its comparison with relational databases. <https://api.semanticscholar.org/CorpusID:15594476>

3. Among the **data models supported by NoSQL databases are key-value, document, column-family and graph**. Every data model offers many methods of data organization and querying and is designed for certain use cases. NoSQL databases include, for instance, Apache Cassandra, MongoDB and Apache HBase. (Nayak, 2013)<sup>21</sup>.

## 2.5 Data Integration

Many sources data are merged into a single viewpoint or format in data integration. It means changing data architecture, resolving contradictions and ensuring that data is consistent and of high quality among many systems.

Generally speaking, **data sources are heterogeneous**, meaning that data often resides in numerous systems — such as databases, files, APIs and streaming platforms — with various schemas, formats and access mechanisms. Managing many data formats, protocols and APIs is necessary to combine data from several unrelated sources. Additionally possible are duplicates, missing figures, contradictory information or disparities in data from different sources. These issues need to be addressed by data integration processes using techniques for data cleansing, deduplication and validation. Field names and data types are examples of dynamic structures that may change with time, as can data schemas. Data integration pipelines need to be able to handle changes to schemas and keep them up to date without affecting processes further down the line.

Data integration often entails taking data out of source systems, converting it to a standard format and utilizing ETL and ELT methods, loading it into a destination system. With Apache Spark and other big data frameworks, ETL operations may be implemented. One such way is to use an API to get data. Scala libraries like Akka HTTP and Play Framework let programmers communicate with RESTful APIs and enable seamless data transmission across web services and cloud platforms. Finally, without needing real data transfer or copying, data virtualization provides a single view of data. Construction of a virtual layer that abstracts the complexity of underlying data sources and provides a single interface for data access and querying is required. As Scala allows for the creation of scalable and effective data access layers, it is ideally suited to implement data virtualization ideas.

### 2.5.1 Data Pipeline Architectures

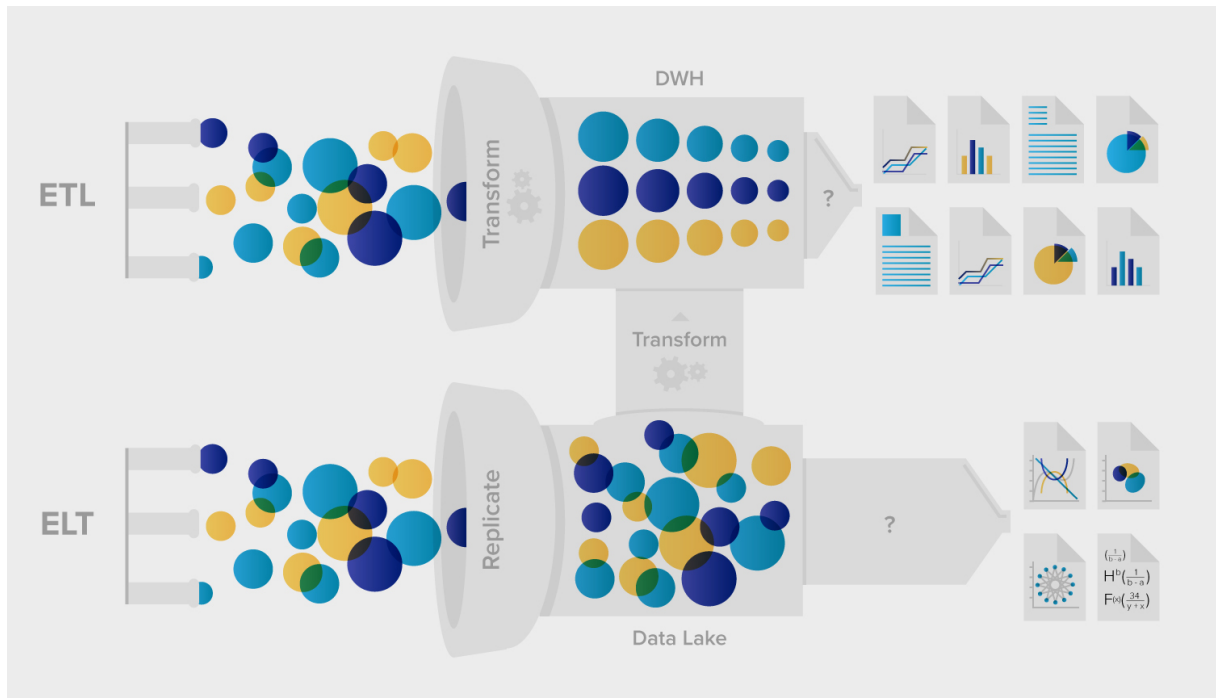
**ETL (Extract, Transform, Load) and ELT (Extract, Load, Transform)** are two fundamental data processing paradigms used in data engineering to extract data from various sources, transform it into a suitable format and load it into target systems such as data warehouses or data lakes (Winn, 2023)<sup>22</sup>.

---

<sup>21</sup>Nayak, A. (2013). Type of nosql databases and its comparison with relational databases. <https://api.semanticscholar.org/CorpusID:15594476>

<sup>22</sup>Winn, J. (2023). Etl vs elt: Understanding the differences and making the right choice. Retrieved June 15, 2024, from <https://www.datacamp.com/blog/etl-vs-elt>

Figure 2.1: ETL and ELT



*Note.* The image depicts a comparison between two data processing approaches. ETL (Extract, Transform, Load) and ELT (Extract, Load, Transform). Both processes start with diverse data sources represented by colored circles; ETL transforms data before loading into a data warehouse (DWH); ELT loads data into a data lake first, then transforms it. The outcomes of both processes are visualized as various types of reports and charts *Creator:* (Winn, 2023)<sup>22</sup>

## 2.5.2 ETL (Extract, Transform, Load)

**ETL** is a process where data is first extracted from source systems, then transformed into a consistent format and finally loaded into the target system (Winn, 2023)<sup>22</sup>.

- In the **extraction phase**, data is gathered from several places during the extraction stage, including databases, files, APIs and streaming services. This phase entails managing many data formats, guaranteeing data consistency and resolving data quality problems; (Winn, 2023)<sup>22</sup>;
- To meet the destination system's schema, the extracted data is cleaned, validated and reorganized during the **transformation stage**. Data type conversions, data enrichment, data aggregation and the application of business rules or logic are all examples of data transformations that are often used (Winn, 2023)<sup>22</sup>;
- **Data is loaded** into the destination system — usually a data warehouse or data lake — after the transformation. This stage takes into account elements such data availability, integrity and volume (Winn, 2023)<sup>22</sup>.

<sup>22</sup>Winn, J. (2023). Etl vs elt: Understanding the differences and making the right choice. Retrieved June 15, 2024, from <https://www.datacamp.com/blog/etl-vs-elt>

When the destination system has limited processing power or the source data needs major modifications, ETL is often the better option. It makes complicated data transformations possible to be carried out prior to data loading, guaranteeing a consistent structure in the target system (Winn, 2023)<sup>22</sup>.

### 2.5.3 ELT (Extract, Load, Transform)

An other method called **ELT** loads the data straight into the target system once it has been retrieved from the source systems; the transformations take place after the loading procedures (Winn, 2023)<sup>22</sup>.

- Like ETL, which gathers data from many sources and loads it into the target system, the **extraction and loading stages are comparable**. Nevertheless, the data is put in ELT in its original format without any previous processing (Winn, 2023)<sup>22</sup>;
- Utilising the processing power and capabilities of the target system, the **transformation phase** of ELT occurs there. Because the target system can effectively manage huge amounts of data, this enables more flexible and scalable data transformations (Winn, 2023)<sup>22</sup>.

Platforms for cloud computing and big data are making ELT more and more prevalent. Because transformations may be implemented on-demand depending on particular needs, it facilitates quicker data loading and more flexible data processing. ELT does, however, also bring difficulties like managing schema change in the destination system and guaranteeing data quality. Data integrity maintenance calls for strong data governance and data validation procedures (Winn, 2023)<sup>22</sup>.

### 2.5.4 Comparison of ETL and ELT

The complexity of data transformations, the processing power of the destination system and the particular needs of the data pipeline are only a few of the considerations when deciding between ETL and ELT. When the destination system is underpowered in processing or when the source data needs intricate changes, ETL is an option. In order to preserve both the quality and the integrity of the data, it guarantees that it is standardized before loading. Conversely, ELT works better with big amounts of data and when the destination system can handle it well enough. It makes data loading quicker and offers transformation application flexibility according to certain use cases (Winn, 2023)<sup>22</sup>.

---

<sup>22</sup>Winn, J. (2023). Etl vs elt: Understanding the differences and making the right choice. Retrieved June 15, 2024, from <https://www.datacamp.com/blog/etl-vs-elt>



## 2.6 Data Security and Governance

**Data security** — refers to the measures and practices implemented to protect data from unauthorized access, breaches, corruption and loss. In the context of data engineering, this involves securing data at rest, in transit and during processing.

**Data governance** — refers to the overall management of data availability, usability, integrity and security. The processes include establishing policies, procedures and standards for data management across the organization.

In today's business world, data has emerged as one of the most precious assets that companies possess. Accurate business insights and regulatory compliance depend on robustly controlling and managing data quality as data volumes increase dramatically across many platforms. Poor quality data may mask true insights and result in poor business choices (Achanta and Boina, 2023)<sup>23</sup>. In order to guarantee high-quality, safe and compliant data assets across the company, data governance offers the overall strategy, rules, standards and procedures. Through this process, regulatory, operational and strategic goals are brought into alignment with data management procedures. Important tasks include **creating data rules, standards and processes; controlling data models, metadata and quality; ensuring clear ownership and stewardship of data; and monitoring data usage and access for compliance and security** (Achanta and Boina, 2023)<sup>23</sup>.

Core to governance is data security. Strong security measures must be implemented inside data platforms by enterprises in light of the increasing cybersecurity risks and stringent data protection laws like **GDPR**. This covers tools for data loss protection, activity monitoring, data encryption and finely tuned access controls. Standardization of data management and categorization guarantees sufficient protection of sensitive data (Achanta and Boina, 2023)<sup>23</sup>. Huge amounts of data go through many phases in data engineering pipelines, from ingestion to consumption. Governance and quality validations are crucially ingrained in data infrastructure and procedures by data engineers. Important areas of concentration include robust data validation during intake, data cleaning and standardization inside ETL processes, data quality checks in analytics and data lineage and cataloging (Achanta and Boina, 2023)<sup>23</sup>.

**Absence of data governance results in data silos, erroneous data definitions, bad metadata and insufficient compliance.** Data exchange is hampered by this, which also damages data trust and raises regulatory concerns. Business, IT and compliance teams must be represented in the well-designed data governance architecture that is implemented. Definition of essential data items, reference data, data quality standards and data problem resolution need the establishment of procedures and responsibilities (Xie et al., 2021)<sup>24</sup>. New possibilities to scale governance are presented by emerging technology. Complicated data matching, quality assessment and anomaly detection may all be automated using machine learning. Tools for managing metadata provide impact analysis, data lineage and auto-classification. Users may find and comprehend data assets by using the business glossaries and data

---

<sup>23</sup>Achanta, A., & Boina, R. (2023). Data governance and quality management in data engineering. *International Journal of Computer Trends and Technology*. <https://api.semanticscholar.org/CorpusID:265699088>

<sup>24</sup>Xie, Q., Zhang, H., Tang, Y.-r., & Lin, M. (2021). Solution ideas and practices for data governance engineering in colleges and universities. *E3S Web of Conferences*. <https://api.semanticscholar.org/CorpusID:235873740>



dictionaries included in data catalogs. Templates and procedures from cloud-based data governance providers help jumpstart projects (Achanta and Boina, 2023)<sup>23</sup>.

## 2.7 Real-World Data Security Challenges and Solutions

The following subchapter is divided into four parts. The first part is a brief definition of the topic and the next three parts are example, challenge and solution, respectively.

**Encryption** — is the process of encoding data so that it can only be decoded with the right encryption key.

- For instance, a financial services organization may utilize AES-256 encryption to safeguard private client data kept in its data warehouse. Data is encrypted in transit by the organization using TLS (Transport Layer Security) as it is transmitted between its on-premises systems and cloud-based analytics platforms.
- The problem is the advancements in quantum computing; there is increasing worry that these machines may be able to crack present encryption systems.
- Organizations looking to future-proof their data security are starting to investigate and use quantum-resistant encryption techniques, often referred to as post-quantum cryptography.

**Access control** — Ensuring that only authorized personnel may access certain data sets requires the implementation of strong access control systems.

- In their data lake, a healthcare provider may, for instance, use role-based access control (RBAC). All patient data may be available to doctors, but billing information is the only information administrative personnel may view. Anonymized data sets may be available to researchers.
- Controlling access permissions across many systems and platforms becomes more difficult as businesses expand and data is dispersed.
- A great solution are Identity and Access Management (IAM) systems that provide centralized management over user authentication and authorization across many systems and cloud platforms are being used by many companies. For access management that is both more granular and aware of the context, some people are also investigating the usage of attribute-based access control, often known as ABAC.

**Data masking and anonymization** — When sensitive data must be utilized for analytics or testing, data masking and anonymization methods are used to preserve data usefulness while safeguarding individual privacy.

---

<sup>23</sup>Achanta, A., & Boina, R. (2023). Data governance and quality management in data engineering. *International Journal of Computer Trends and Technology*. <https://api.semanticscholar.org/CorpusID:265699088>

- A retail corporation may, for instance, utilize data masking methods to provide its marketing analytics team with consumer transaction data. The statistical features of the original dataset are preserved, while realistic but fictitious data replaces personally identifiable information (PII) such as names and addresses.
- Finding the appropriate equilibrium between the value of data and the preservation of privacy may be challenging, particularly when applied to more sophisticated re-identification systems.
- Differential privacy approaches are becoming more popular among firms. These techniques include the addition of noise to datasets that have been properly calibrated in order to give robust privacy assurances as well as the ability to do relevant analysis.

## **Data catalog and metadata management**

**Metadata** — data that describes other data

**Data catalog** — centralized database that organizes and saves metadata on the data assets of an organization. It serves as an inventory system for data, providing information about data sources, schemas, relationships and usage.

**Metadata management** — is the process of creating, storing organizing and maintaining metadata.

- An international company may, for instance, put in place a data catalog system that offers a searchable inventory of all data assets within the company. Information regarding the history of the data, its owners, measures of data quality and use rules is included in this catalog.
- Keeping an accurate and current data catalog is harder as data quantities and diversity rise.
- For automated data discovery and categorization, several companies are using machine learning and AI technology. These applications can identify data types, search data repositories automatically and update in real time the data catalog.

**Data quality management** — refers to the set of practices, processes and technologies used to measure, improve and maintain the quality of data within an organization

- An online retailer may, for instance, include automatic data quality checks into their ETL processes. Among these checks might include validating that product pricing are within an expected range, verifying that client email addresses are formatted correctly and flagging any unusual spikes in order volumes that could be a sign of data problems.
- As data sources become more diverse and data volumes increase, maintaining consistent data quality becomes more complex.
- Numerous companies are using sophisticated data observability systems. These systems may automatically identify and notify on abnormalities across many data quality parameters and utilize machine learning to set baselines for data quality.

**Regulatory compliance** — refers to an organization’s adherence to laws, regulations, guidelines and specifications relevant to its business processes. In data engineering and governance, refers to making sure that data management procedures, security protocols and privacy safeguards satisfy the criteria established by different regulatory agencies and business associations.

- A worldwide bank must, for instance, abide by a number of laws, such as the California Consumer Protection Act (CCPA) and industry-specific rules like PCI DSS for payment card data. They implement compliance management system that monitors consent, data flows and offers audit trails for every data processing and accessing operation.
- New laws are established and current ones are amended on a regular basis, thus the regulatory environment is ever changing.
- To keep informed about regulatory changes, evaluate their effect and automate compliance procedures, many companies are using RegTech (Regulatory Technology) solutions that leverage AI and machine learning.

#### **Data ethics and responsible AI**

- An AI ethics board that evaluates all AI projects is implemented by a big tech corporation that creates AI models for content recommendation. They also put into place methods to identify and reduce bias in their model outputs and training data.
- It may be difficult to define and apply moral principles for AI and data usage, particularly when addressing complicated situations or competing agendas.
- A few groups are embracing official guidelines for ethical AI, notably the EU’s Ethics Guidelines for Trustworthy AI. When it comes to addressing ethical problems in the creation and deployment of artificial intelligence, these frameworks provide an organized approach.

## **2.8 Data Serving**

Operating machine learning requires data pipelines to automate data flow from source to model training to inference. Because of shortcomings in the tools, the pipeline must smoothly combine stream processing with model serving for real-time applications. Particularly difficult is serving machine learning models on streaming data when the models must be updated often via online learning. Essential MLOps procedures include model versioning, monitoring, auditing and zero-downtime upgrades. In production, dynamic learning is made possible by scalable systems that use model serving frameworks and stream processing platforms like Kafka. The end-to-end latency and throughput are greatly influenced by the model serving setup’s performance. Understanding the subtle variations between different streaming and serving technologies requires comparative benchmarking. For certain tasks, GPU-based hardware acceleration may be very beneficial. Semantic technologies are being used more and more to make

building and using ML pipelines easier. Data science may be efficiently used by people who aren't ML specialists by capturing domain knowledge and ML best practices in ontologies. These ontologies can then be accessible via intelligent modules and intuitive interfaces. This has use in industrial environments (Cédola et al., 2021)<sup>26</sup>.

## 2.9 Monitoring and Maintenance

Machine learning systems' foundation are data pipelines, which automate the flow of data from intake to model training and inference. The provision of machine learning models with timely and high-quality data to train from is made possible by a data pipeline that has been well-architected. It is important to monitor data flows so that problems are caught early and the system stays healthy generally. This includes tracking metrics on pipeline performance, data distribution, schema changes and data volume. Anomaly detection methods can be used to find quick changes and let engineers know about them. Considerable work is put into feature engineering and turning unprocessed data into features for machine learning models during pipeline construction. Combining transaction data, analyzing text sentiment, computing time-series statistics, etc. How well a model works depends a lot on how good its features are. Pipelines must be adaptable since data arrives in numerous forms from many sources. When dealing with unstructured data such as video, it may be necessary to use specific procedures. One example of this would be filtering video frames based on their quality and content in order to generate a training dataset that is representative. **Scalable data pipeline deployment has never been simpler thanks to cloud platforms.** So that data engineers may concentrate on pipeline logic, managed services take care of infrastructure provisioning and orchestration. Kubeflow is one tool that helps connect pipeline elements. Additionally, under active development is data versioning. Debugging and reproducibility of a certain model depend on the ability to monitor the version of data used to train it. Data lineage monitoring and pipeline metadata collecting make this possible. In conclusion, one of the most important applications of machine learning pipelines is the use of predictive maintenance in sectors such as manufacturing, energy and transportation. By analyzing sensor data and event records, models can anticipate equipment failures in advance, enabling the scheduling of restorations proactively. This enhances reliability and reduces costs. In conclusion, data pipeline seamless functioning depends on monitoring and maintenance, which in turn makes data-driven projects successful in many sectors. To keep the data flowing, proactive maintenance, careful monitoring and sound system design are all necessary (Roychowdhury and Sato, 2021)<sup>25</sup>.

---

<sup>26</sup>Cédola, A. P., Rossini, R., Bosi, I., & Conzon, D. (2021). Feature engineering and machine learning modelling for predictive maintenance based on production and stop events. <https://api.semanticscholar.org/CorpusID:244727976>

<sup>25</sup>Roychowdhury, S., & Sato, J. Y. (2021). Video-data pipelines for machine learning applications

## Chapter 3

# Parallel and distributed data processing

Unlocking the potential of big data depends critically on parallel and **distributed processing (DP)**. Classical centralized methods become unsustainable and ineffective as data quantities grow exponentially. Big data workloads may be split up into smaller, more manageable jobs that can be carried out concurrently on many processing units by using **parallel processing (PP)**. Because of parallelism, huge speedups are possible and it is possible to handle enormous datasets in a reasonable amount of time.

By distributing data and calculations over a cluster of computers, **DP** enhances parallelism. Both scalability and fault tolerance are provided by this configuration. To handle the increasing processing burden, more computers may be added to the cluster as the data volume rises. Because they duplicate data over many nodes, distributed file systems — like the **Hadoop Distributed File System (HDFS)** — ensure data availability and dependability.

By providing abstractions and **application programming interfaces (API)** that make it simpler to construct distributed systems, **Hadoop** and **Spark** have brought about a dramatic shift in the way that large data is handled. Using **MapReduce** functions, developers may specify calculations in **Hadoop** that are automatically parallelized and carried out on the cluster. By adding in memory processing capabilities and a large number of APIs for data manipulation, machine learning and graph processing, Spark expands on this idea.

**PP**, **DP** and **cloud infrastructure** working together have given big data analytics new opportunities. Elastic resources on cloud platforms may be dynamically provided and scaled according to the demands of the workload. Organizations may now handle and analyze enormous datasets without having to make large upfront infrastructure expenditures.

But efficiently using **PP** and **DP** for big data for thorough analysis of resource management, task scheduling and data partitioning. The optimization of the allocation of tasks to available resources, the minimization of data transportation and the guaranteeing of load balance are all important functions that are performed by **scheduling algorithms**. Further sophisticated methods, like adaptive resource allocation and data locality-aware scheduling, may raise large **DP's** effectiveness and performance even further.

### 3.1 Fundamentals of Parallel and Distributed Processing for Big Data

Big data workloads have outgrown conventional sequential processing methods as data workloads have grown exponentially in recent years. A key **paradigm for effectively analyzing and obtaining value from enormous datasets is parallel and distributed processing**. "Why parallel and distributed processing is essential for big data" is the topic this part attempts to address.

Data sets involved in **big data** often surpass the storage and processing capacities of a single system by a significant margin, measuring in petabytes or exabytes. The scalability that is required to manage these enormous amounts of data may be achieved via the use of parallel and distributed architectures, which enable the data and computation to be dispersed among clusters of hundreds or thousands of commodity servers (Fadiya and Sari, 2018)<sup>28</sup>.

Distributed systems have the potential to significantly decrease processing times in comparison to sequential alternatives. This is achieved by **distributing data and tasks over a large number of nodes** that are capable of working in parallel. This offers significantly quicker insights and near-real-time analytics on large datasets. (Natesan et al., 2023)<sup>27</sup>.

When compared to the use of costly supercomputers or specialist hardware, the utilization of distributed computing on commodity hardware clusters is much more **cost-effective** for the processing of large amounts of data (Fadiya and Sari, 2018)<sup>28</sup>.

Large dataset processing over lengthy periods of time requires distributed frameworks to be able to **gracefully accept individual node failures** without impacting the overall workload (Wingerath et al., 2016)<sup>29</sup>. Few points sum up the principles of distributed and parallel computing.

- Bigger datasets are **divided into smaller chunks** that may be handled concurrently and independently by many nodes. Typical methods include round-robin, hash and range partitioning;
- Often following patterns like **MapReduce, bulk synchronous parallel (BSP)** and **dataflow/stream processing**, algorithms are built to work on partitioned data in parallel.
- Specialized file systems, such as the **Hadoop Distributed File System (HDFS)**, provide a single view and enable data to be stored across a cluster.
- Resources (CPU, memory, etc.) are allocated to several simultaneous activities via **resource management** frameworks such as YARN;

---

<sup>28</sup>Fadiya, S., & Sari, A. (2018). The importance of big data technology. *International Journal of Engineering & Technology*. <https://api.semanticscholar.org/CorpusID:199019074>

<sup>27</sup>Natesan, P., E, S. V., Mathivanan, S. K., Venkatesan, M., Jayagopal, P., & Allayear, S. M. (2023). A distributed framework for predictive analytics using big data and mapreduce parallel programming. *Mathematical Problems in Engineering*. <https://api.semanticscholar.org/CorpusID:256524905>

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>

- Parallel jobs require **coordination and synchronization** techniques, as does shared state across distributed nodes.

## 3.2 Distributed Processing Frameworks for Big Data

Among the first programming methods for distributed large data processing was **MapReduce**, made famous by Google. With the release of an open-source MapReduce implementation along with HDFS, Apache Hadoop emerged as the industry standard for batch processing large data (Fadiya and Sari, 2018)<sup>28</sup>.

**MapReduce and Hadoop** key characteristics are: Simple programming model with Map and Reduce phases; Automatic parallelization and distribution; Fault-tolerance through data replication and speculative execution; Scalability to thousands of nodes; However, MapReduce has limitations for iterative algorithms and real time processing. **Apache Spark**: in-memory processing for faster performance; Rich APIs in multiple languages; Support for batch, interactive and stream processing; Advanced analytics capabilities (machine learning, graph processing, etc.); Spark's Resilient Distributed Datasets (RDDs) provide fault-tolerant distributed memory abstraction. **Apache Storm**: Low-latency stream processing; At-least-once processing semantics; Native streaming (not micro-batch). **Apache Flink**: Both batch and stream processing; Exactly-once processing semantics; Lower latency than Spark Streaming. **Apache Samza**: Built for high-throughput event processing; Tightly integrated with Apache Kafka (Wingerath et al., 2016)<sup>29</sup>.

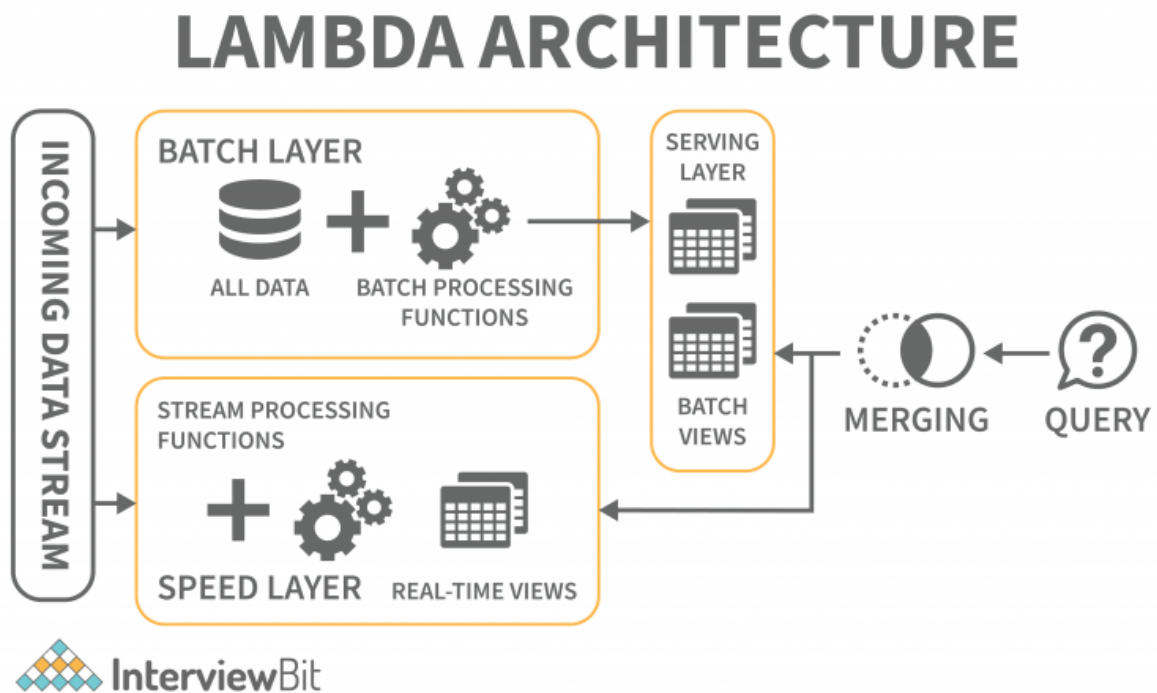
## 3.3 Architectural Patterns for Big Data Processing

The **Lambda Architecture**, proposed by Nathan Marz, combines batch and stream processing to balance latency, throughput and fault-tolerance. It consists of three layers: **Batch Layer**: Periodically processes all historical data using a system like **Hadoop**; **Speed Layer**: Processes real time data streams for low-latency views; **Serving Layer**: Responds to queries by merging batch and real time views. This architecture provides comprehensive and accurate batch analytics along with approximate real time results (Wingerath et al., 2016)<sup>29</sup>.

<sup>28</sup>Fadiya, S., & Sari, A. (2018). The importance of big data technology. *International Journal of Engineering & Technology*. <https://api.semanticscholar.org/CorpusID:199019074>

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>

Figure 3.1: Lambda Architecture



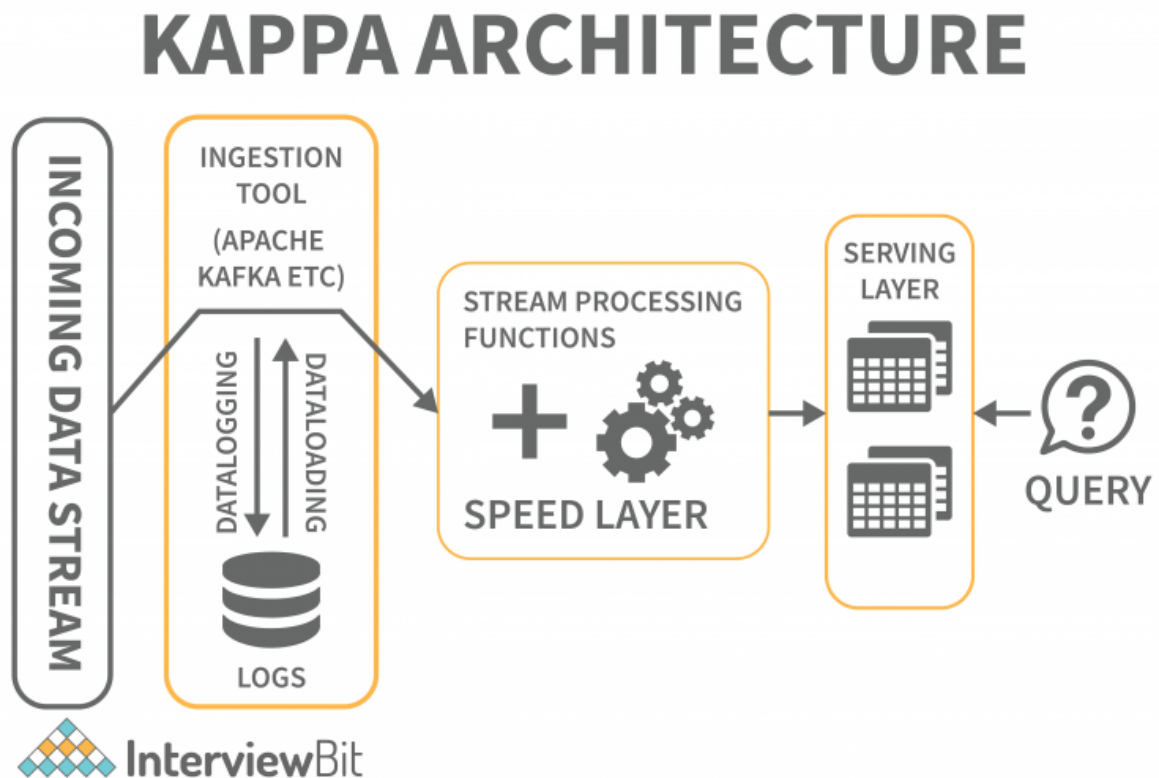
*Note.* The image depicts the Lambda architecture, a data processing framework designed for handling big data in real-time applications. **Batch Layer** — Processes all data and performs batch processing functions; **Speed Layer** — Handles stream processing and real-time views; **Serving Layer** — Provides batch views and merges data for queries; **Incoming Data Stream** — Feeds into both batch and speed layers; **Merging** — Combines batch and real-time processed data. *Creator:* (InterviewBit, 2022)<sup>39</sup>

The **Kappa Architecture**, introduced by Jay Kreps, simplifies the **Lambda Architecture** by treating both real time and historical data as streams. It uses a single stream processing engine capable of high-throughput reprocessing of historical data. **Key components:** Scalable message queue (e.g., Apache Kafka); Stream processor (e.g., Apache Samza); Serving database for processed results. This approach reduces complexity but requires a very robust stream processing system (Wingerath et al., 2016)<sup>29</sup>.

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>



Figure 3.2: Kappa Architecture



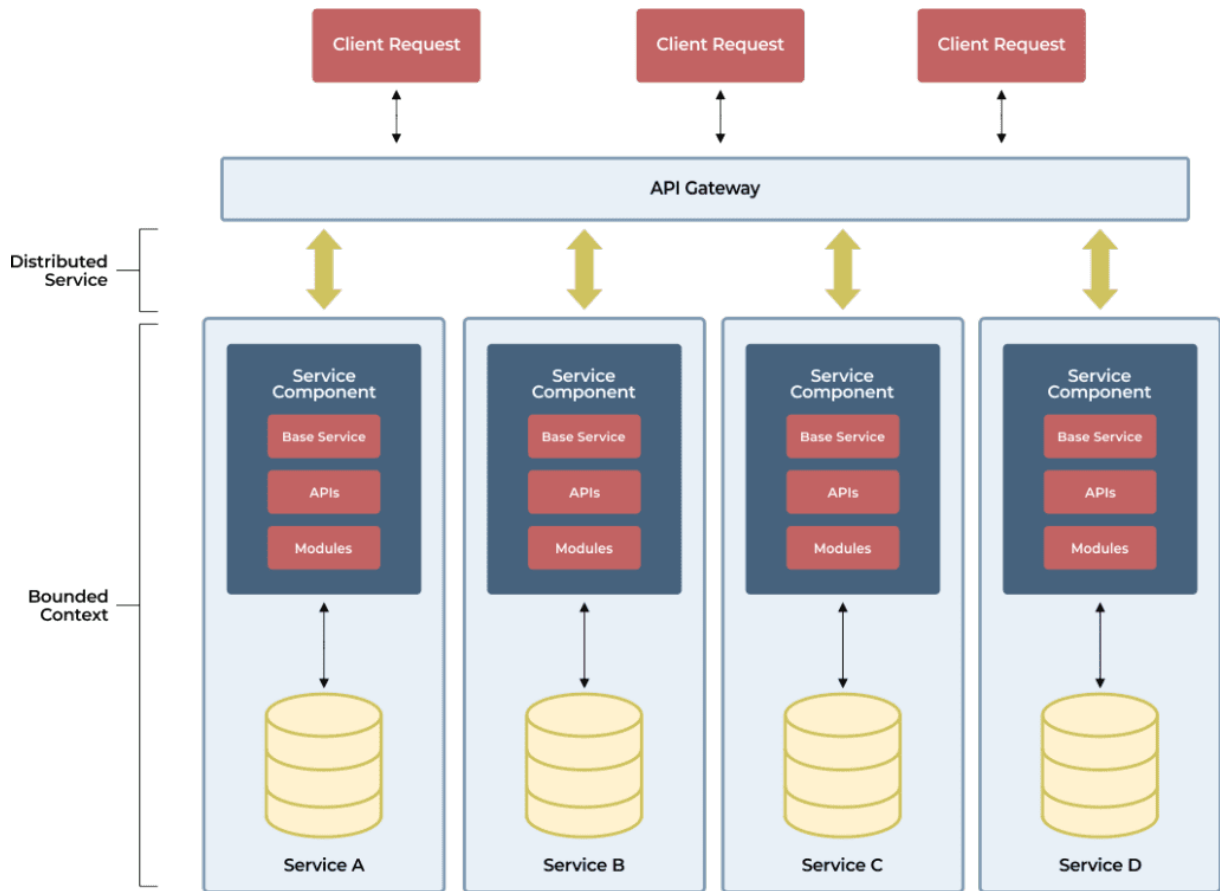
*Note.* The image depicts the Kappa Architecture, which is a data processing framework designed for handling real-time data streams. **Incoming Data Stream** — The starting point of the data flow; **Ingestion Tool** — Utilizes technologies like Apache Kafka for data intake; **Stream Processing Functions** — Core data processing layer; **Speed Layer** — Handles real-time data processing; **Serving Layer** — Presents processed data for querying; **Query Interface** — Allows users to interact with the processed data; *Creator:* (InterviewBit, 2022)<sup>39</sup>

**Microservices architectures** decompose big data applications into smaller, loosely-coupled services that can be developed, deployed and scaled independently. This approach can improve agility and scalability for complex big data systems (Wingerath et al., 2016)<sup>29</sup>, (InterviewBit, 2022)<sup>39</sup>

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>

<sup>39</sup>InterviewBit. (2022). Big data architecture – detailed explanation. <https://www.interviewbit.com/blog/big-data-architecture/>

Figure 3.3: Microservice Architecture Style Structure



*Note.* The image depicts a microservices architecture diagram. At the top of the diagram, there are three "Client Request" boxes, representing incoming requests from clients. Below the client requests is an "API Gateway" layer, which serves as the entry point for all client requests. The main body of the diagram shows **four separate services** (Service A, B, C, and D), representing a distributed service architecture. *Creator:* (Jain, 2020)<sup>40</sup>

### 3.4 Parallel Collections

A tool developed to facilitate parallel programming and enhance multicore system performance is the **parallel collections (PC)** framework in Scala. Developers may take use of the potential of parallel computing without having to deal with low-level concurrency specifics thanks to this high-level abstraction for parallel data processing (Prokopec et al., 2024)<sup>30</sup>.

In Scala, **PCs** run by splitting up data into pieces and running many threads at once to process each chunk. Through the use of this divide-and-conquer strategy, it is possible to make effective use of the computer resources that are available. (Prokopec et al., 2024)<sup>30</sup>.

<sup>30</sup>Prokopec, A., Miller, H., & Contributors. (2024). Parallel collections overview. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/overview.html>

Table 3.1: Parallel vs. sequential processing

```
val result = (1 to 1000000).filter(_ % 2 == 0).map(_ * 2).sum // Sequential processing
val parallelResult = (1 to 1000000).par.filter(_ % 2 == 0).map(_ * 2).sum // Parallel processing
```

*Note.* In this example, the parallel version (par) can potentially execute much faster on multi-core systems, especially for large datasets. *Creator.* Author’s own work.

Scala’s **PCs** framework provides parallel counterparts for most of the standard sequential collections including: **ParArray**, **ParVector**, **ParRange**, **ParSet** and **ParMap**. Each of these collections supports parallel operations while maintaining the semantics of their sequential counterparts (Prokopec et al., 2024)<sup>30</sup>.

Though **PCs** may boost performance significantly, it is important to know when and how to utilize them. **They work well for:** Operations that benefit from parallelization more are those that need a lot of processing for each element; The overhead of parallelization is more justified for larger collections; Parallel processing works best for jobs when each component can be handled separately without interfering with others. **In terms of potential pitfalls:** Small collections may find that the costs of parallelization exceeds the advantages; Parallelization may not benefit I/O-intensive tasks; Due to the fact that some processes are non-associative, it is possible that when they are parallelized, they will generate different outcomes compared to their sequential counterparts (Prokopec et al., 2024)<sup>30</sup>.

Table 3.2: Non-associative operation

```
val numbers = (1 to 1000).toList
println(numbers.reduce(_ - _)) // Sequential - predictable result
println(numbers.par.reduce(_ - _)) // Parallel - potentially unpredictable result
```

*Note.* In this example, subtraction is non-associative, and the parallel version may produce different results in different runs. *Creator.* Author’s own work.

### 3.4.1 Task Support (TS)

By use of the idea of **Task Support**, the Scala PCs framework enables **customization of task scheduling and execution**. Generally speaking, it is a part that offers fine-grained management of parallel operation execution. Customizing the execution environment for parallel operations is made possible by its function as an abstraction layer between the parallel collections and the underlying thread pool technology (Contributors, 2024)<sup>31</sup>.

Scala’s **PCs** by default employ a **ForkJoinTaskSupport** implementation that is derived from the Fork/Join framework seen in Java. For most general-purpose parallel processing requirements, this default solution works well. But Scala offers the ability to alter this behavior (Contributors, 2024)<sup>31</sup>.

<sup>30</sup>Prokopec, A., Miller, H., & Contributors. (2024). Parallel collections overview. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/overview.html>

<sup>31</sup>Contributors. (2024). Task support. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/configuration.html#task-support>

Table 3.3: Non-associative operation

```
import scala.collection.parallel.TaskSupport
import scala.collection.parallel.ForkJoinTaskSupport
import java.util.concurrent.ForkJoinPool
val customParSeq = (1 to 10000).par
customParSeq.tasksupport = new ForkJoinTaskSupport(new ForkJoinPool(4))
```

*Note.* In this example, a custom **ForkJoinTaskSupport** is created with a specific **ForkJoinPool** that limits the parallelism to four threads. This level of control is particularly useful in scenarios where it is needed to manage system resources carefully or when working in environments with specific threading requirements. *Creator:* Author's own work.

- **TS** decides how to divide, plan and control jobs among the available threads or processors;
- Depending on particular application requirements, it enables to choose or apply several task execution techniques;
- Through **TS**, the level of parallelism and resource utilization in their applications can be controlled (Contributors, 2024)<sup>31</sup>.

### 3.4.2 Performance Implications of Task Support

- Various **TS** implementations may modify the way work is distributed across threads, therefore affecting the ratio of parallelism overhead to efficient use of resources;
- Particularly for CPU-bound jobs, the frequency of context transitions may have an influence on overall performance depending on the thread pool design;
- **TS** setup done correctly guarantees best use of the system resources, avoiding CPU core oversubscription or underutilization (Contributors, 2024)<sup>31</sup>.

Table 3.4: Performance tuning with Task Support

```
import scala.collection.parallel.ForkJoinTaskSupport
import java.util.concurrent.ForkJoinPool
def parallelSum(numbers: Seq[Int], parallelism: Int): Int = {
  val par = numbers.par
  par.tasksupport = new ForkJoinTaskSupport(new ForkJoinPool(parallelism))
  par.sum}
val result = parallelSum((1 to 1000000), Runtime.getRuntime.availableProcessors())
```

*Note.* In this example, the parallelism is set to match the number of available processors, which can lead to more efficient execution on the given hardware *Creator:* Author's own work.

When working with Task Support a few thing should be considered.

<sup>31</sup>Contributors. (2024). Task support. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/configuration.html#task-support>

- What is the application nature? CPU-bound vs. I/O-bound tasks may benefit from different **TS** configurations;
- To prevent overloading the level of parallelism has to be aligned with available resources;
- The size and complexity of the data being processed should determine TS strategy.
- Lastly, benchmarking. Different **TS** configurations should be tested to find the optimal setup for specific use case (Contributors, 2024)<sup>31</sup>.

### 3.5 Futures and Async Programming

A **Future** represents a value that may not yet be available but will be at some point in the future. Handling asynchronous computations — like I/O operations or lengthy computations — without interrupting the primary execution thread depends on this notion. Non-blocking execution made possible by Futures also enables effective use of available resources. Finally, complicated asynchronous operations are made possible by the combination and transformation of many **Futures** (Haller et al., 2024)<sup>32</sup>.

Table 3.5: Futures

```
import scala.concurrent.Future
import scala.concurrent.ExecutionContext.Implicits.global
def fetchData(url: String): Future[String] = Future {
  Thread.sleep(1000) // Simulating a network call
  s"Data from $url"}
val dataFuture: Future[String] = fetchData("http://example.com/data")
```

*Note.* In this example, **fetchData** returns a **Future[String]**, allowing the program to continue execution while the data is being fetched asynchronously. *Creator:* Author's own work.

Because Scala's **asnc/await** feature makes working with **Futures** more natural, asynchronous programming seems more sequential and reasonable. This helps especially with complicated jobs that need for coordination of many asynchronous actions. (Haller et al., 2024)<sup>32</sup>.

### 3.6 Akka Framework

Akka is a toolkit and runtime for building highly concurrent, distributed and resilient message-driven applications on the JVM. It provides a higher level of abstraction for writing concurrent and distributed systems.

<sup>31</sup>Contributors. (2024). Task support. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/configuration.html#task-support>

<sup>32</sup>Haller, P., Prokopec, A., Miller, H., Klang, V., Kuhn, R., Jovanovic, V., & Contributors. (2024). Futures and promises. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/core/futures.html>

Table 3.6: Using Async/Await

```
import scala.concurrent.Future
import scala.concurrent.ExecutionContext.Implicits.global
import scala.async.Async.{async, await}
def processData(data: String): Future[Int] = Future {data.length} // Simulating data processing
val result: Future[Int] = async {val data = await(fetchData("http://example.com/data"))
  await(processData(data))}
```

*Note.* In this example, **async** creates a block that can contain **await** calls, which pause execution until the **Future** completes, without blocking the thread. This approach simplifies the handling of multiple dependent asynchronous operations. *Creator:* Author's own work.

Table 3.7: Error handling

```
def fallbackData: Future[String] = Future.successful("Fallback Data")
val robustDataFetch: Future[String] = fetchData("http://example.com/data")
  .recover { case _: Exception => "Default Data" }.fallbackTo(fallbackData)
```

*Note.* In this example, **recover** handles specific exceptions and provides a default value, and **fallbackTo** tries an alternative **Future** if the first one fails. This approach ensures that data pipelines can gracefully handle failures and continue processing with default or fallback data. *Creator:* Author's own work.

Table 3.8: Composition

```
def processMultipleDataSources(): Future[List[Int]] = {
  val dataSources = List("http://example.com/data1", "http://example.com/data2")
  Future.traverse(dataSources) { url =>
    for {
      data <- fetchData(url)
      result <- processData(data)
    } yield result}}}
```

*Note.* This example demonstrates how multiple data sources can be processed in parallel. Futures can be composed to create complex workflows; Sequencing using **flatMap** or for-comprehensions or using **Future.traverse**. *Creator:* Author's own work.

### 3.6.1 Actor Model

Programming idea and tool the **Akka Actor Model** was developed to make concurrent and distributed system development easier. For creating **parallel, distributed and failure-resistant applications**, it provides a higher degree of abstraction. Handling systems with components interacting in real time is particularly beneficial using this method, as in large-scale data processing jobs, Internet of Things applications, or responsive web services (Kuhn et al., 2017)<sup>33</sup>.

Fundamentally, the **Actor Model** centers on the notion of **"actors"** as the building blocks of computing. **Lightweight and self-contained, an actor captures state and behavior.** Actors interact only via message passing,

<sup>33</sup>Kuhn, R., Hanafée, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications

which promotes component flexibility and aids in concurrent system complexity control (Kuhn et al., 2017)<sup>33</sup>.

In the Akka implementation of the Actor Model, each actor possesses characteristics;

1. An actor has its own private state that cannot be directly accessed or altered by other actors. This encapsulation aids in preventing concurrency issues such, as race conditions;
2. Actors process messages one at a time, in a single-threaded manner within the actor. This characteristic eliminates the need for complex locking mechanisms within an actor's logic;
3. When an actor receives a message it can change its state, communicate with actors, spawn new actors or decide how to process the next message;
4. Actors, in the system, are designed to work whether they are situated locally (within the JVM) or remotely (on a different machine within the network).

An important advantage of the **Actor Model** is its approach to handling errors. In Akka actors are structured in a hierarchy known as a "**supervisor hierarchy**". If an actor encounters an error (such as throwing an exception) its supervisor. The actor that created it. Determines how to address the issue. This response could involve restarting the actor, halting its operation or escalating the problem to a higher level supervisor. This hierarchical supervision system enables management of failures. Facilitates graceful recovery from them (Kuhn et al., 2017)<sup>33</sup>.

Leveraging the Actor Model also provides a means of application scalability. Distribution of actors among computers in a cluster is quite simple since actors communicate via messages. Clustering capabilities included into Akka allow actors to be dispersed across network nodes while addressing distributed computing issues such network communication and failure management (Kuhn et al., 2017)<sup>33</sup>.

A well-known characteristic of the Akka Actor Model is its ability to support reactive programming ideas. Akka actors are designed to stay responsive, under loads recover from failures, scale based on demand and rely on asynchronous message passing for communication. With its characteristics, Akka is an excellent option for creating reactive systems that can manage concurrency and remain responsive under changing loads (Kuhn et al., 2017)<sup>33</sup>.

When building data processing pipelines in the area of data engineering, the Akka Actor Model shows useful. Actors are components of a data pipeline that stand for certain processing or transformation operations. These pipeline stages are composed more simply with the message passing technique, and the fault tolerance features that are incorporated improve the reliability of data processing processes (Kuhn et al., 2017)<sup>33</sup>.

As an example, some actors may be responsible for gathering data from various sources, while others would handle cleaning and transforming the data. Still others would be tasked with aggregating and evaluating the findings. For better performance, the actor model allows these procedures to be distributed over a machine cluster and parallelized (Kuhn et al., 2017)<sup>33</sup>.

---

<sup>33</sup>Kuhn, R., Hanafee, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications

Table 3.9: Basic Actor Structure in Akka

```
import akka.actor.{Actor, ActorSystem, Props}
class DataProcessorActor extends Actor {var processedCount = 0
  def receive: Receive = {case data: String =>
    processedCount += 1
    println(s"Processing: $data (Total processed: $processedCount)")
    case "getCount" => sender() ! processedCount
    case _          => println("Unknown message")}}
val system = ActorSystem("DataEngineeringSystem")
val dataProcessor = system.actorOf(Props[DataProcessorActor], "dataProcessor")
dataProcessor ! "Sample data"
dataProcessor ! "getCount"
```

*Note.* In this example, **DataProcessorActor** encapsulates its state (**processedCount**) and defines its behavior in the **receive** method. It can process data and respond to queries about its internal state, all through message passing . *Creator:* Author's own work.

Table 3.10: Actor Hierarchies and Supervision

```
class SupervisorActor extends Actor {
  val child1 = context.actorOf(Props[ChildActor], "child1")
  val child2 = context.actorOf(Props[ChildActor], "child2")
  def receive: Receive = { case msg => child1 ! msg // Forward message to child1}
  override val supervisorStrategy = OneForOneStrategy() {
    case _: ArithmeticException => Restart
    case _: Exception           => Stop}}

```

*Note.* In this example, Akka implements a hierarchical structure for actors, known as the **actor hierarchy**. Hierarchy allows for structured error handling and fault tolerance. The supervisor can decide how to handle failures in its child actors, such as restarting them or stopping them . *Creator:* Author's own work.

### 3.6.2 Actors Advantage in Data Engineering

Table 3.11: Actors Scalability and Performance

```
val system = ActorSystem("DataProcessingSystem")
val router = system.actorOf(
  RoundRobinPool(10).props(Props[DataProcessorActor]),
  "dataProcessorRouter")
(1 to 1000).foreach { i =>
  router ! s"Data chunk $i"}
```

*Note.* This example creates a router with 10 instances of **DataProcessorActor**, distributing the workload across multiple actors for parallel processing. Actors are lightweight and can be created in large numbers, allowing for fine-grained parallelism in data processing tasks. *Creator:* Author's own work.



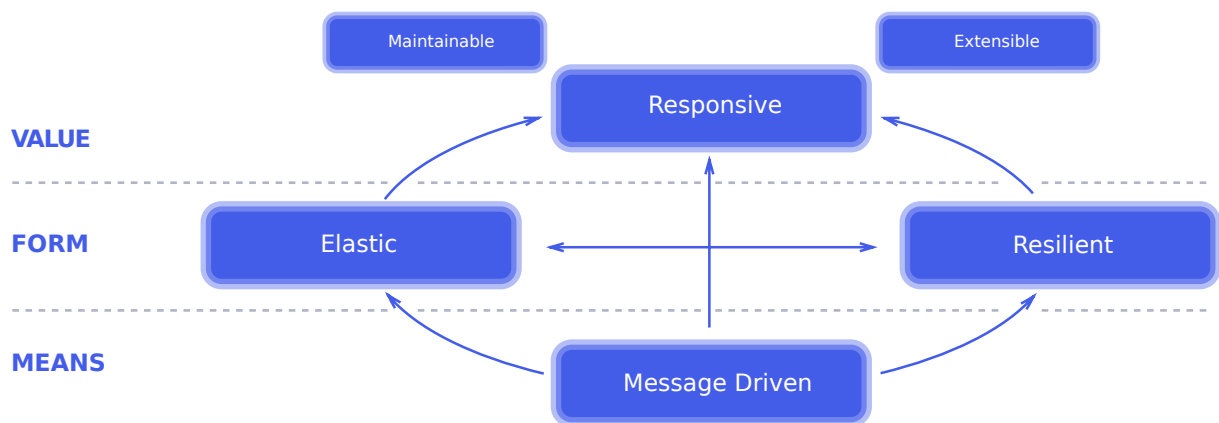
Table 3.12: Actors Scalability and Performance

```
class AggregatorActor extends Actor {
  var sum = 0
  var count = 0
  def receive: Receive = {
    case value: Int =>
      sum += value
      count += 1
    case "getAverage" =>
      sender() ! (if (count > 0) sum.toDouble / count else 0)}
}
```

*Note.* In this example, actor maintains a **running average**, demonstrating how stateful computations can be encapsulated within an actor. Actors can maintain state between messages, useful for stateful data processing operations. *Creator.* Author's own work.

### 3.6.3 Reactive Streams

Figure 3.4: Reactive Streams



*Note.* This diagram illustrates the key principles of Reactive Systems as outlined in the **Reactive Manifesto**. The diagram is divided into three layers: Value, Form, and Means. It showcases six core concepts of Reactive Systems. The layout emphasizes how these concepts interact and support each other. *Creator.* (Lightbend, 2024)<sup>34</sup>

Standard Reactive Streams handle asynchronous stream processing under non-blocking demand. Within the field of distributed and parallel data processing, this technology provides a means of handling enormous amounts of data across several computers. It shows to be useful in data engineering situations where processing pipelines have to control throughput real-time data streams while maintaining system stability and responsiveness (Kuhn et al., 2017)<sup>33</sup>.

At its essence, Reactive Streams establishes a framework of interfaces and rules for facilitating asynchronous stream processing among components within a distributed system. The standard outlines four interfaces; Publisher,

<sup>33</sup>Kuhn, R., Hanafee, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications

Subscriber, Subscription and Processor. These interfaces collaborate to facilitate the flow of data throughout a processing pipeline while effectively managing back pressure to sustain system stability amidst varying workloads (Kuhn et al., 2017)<sup>33</sup>.

Reactive Streams offers various benefits in the field of distributed and parallel data processing; It enables non-blocking asynchronous data processing—a critical feature in distributed systems where components may process data at different speeds. Through asynchronous processing, quicker components are able to function without being delayed by slower ones, which maximizes the efficiency of the system (Kuhn et al., 2017)<sup>33</sup>.

Its integrated system to manage back pressure is one of its main advantages. Different phases of a distributed data processing pipeline may be able to handle jobs to differing degrees. Back pressure is the ability of components to tell upstream components about their processing capabilities, therefore preventing quicker data producers from overpowering slower data users. This contributes to memory-related problems in distributed environments being avoided and system stability being maintained. (Kuhn et al., 2017)<sup>33</sup>.

Reactive Streams supports constructing data processing pipelines modularly. Following the Reactive Streams interfaces, each step of the pipeline may be constructed as reusable components. This allows for more flexibility. Because of its modularity, which is in line with the ideas of functional programming, complex distributed data flows may be created using simple, clearly defined building components (Kuhn et al., 2017)<sup>33</sup>.

Because Reactive Streams are asynchronous and can control back pressure, they are perfect for building scalable distributed data processing systems. More processing nodes may be added into the system as the amount of data grows, and Reactive Streams manages the data flow between them (Kuhn et al., 2017)<sup>33</sup>.

While not included in the Reactive Streams definition directly, implementations often provide capabilities to manage failures in distributed settings. A known implementation of Reactive Streams, Akka Streams, for example, provides features for error handling and recovery in distributed stream processing situations (Kuhn et al., 2017)<sup>33</sup>.

Reactive streams may be used in reality for many different facets of distributed and parallel data processing; building scalable, robust data ingestion pipelines that can handle large volumes of real-time data streams from many sources is one such uses (Kuhn et al., 2017)<sup>33</sup>.

Use of a back pressure mechanism prevents the ingestion process from overpowering subsequent processing phases. Reactive Streams processors, each specialized to a certain transformation aspect, may be used to break down complex data transformation algorithms. Parallel processing is made possible by the dispersal of these processors across nodes. Reactive Streams may be effectively used by real time analytics pipelines to process and analyze data as it moves through the system. Analytical jobs are made easier to complete by Reactive Streams' asynchronous design. Reactive Streams help to control the data flow while exporting processed data to systems so as to avoid overwhelming the receiving systems (Kuhn et al., 2017)<sup>33</sup>.

Reactive Streams standard is followed by many libraries and frameworks in the Scala ecosystem, which makes

---

<sup>33</sup>Kuhn, R., Hanafee, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications

it easier to include these ideas into data engineering projects. As one example, Akka Streams provides a high-level DSL for building Reactive Streams compliant stream processing pipelines. With its smooth integration with the Akkas actor paradigm, distributed and reactive data processing systems may be built (Kuhn et al., 2017)<sup>33</sup>.

Table 3.13: Asynchronous Processing

```
import akka.stream.scaladsl._
import scala.concurrent.Future
def processData(data: String): Future[String] = Future {
  // Simulating asynchronous processing
  Thread.sleep(100)
  data.toUpperCase}
val source: Source[String, NotUsed] = Source(List("data1", "data2", "data3"))
val flow: Flow[String, String, NotUsed] = Flow[String].mapAsync(4)(processData)
val sink: Sink[String, Future[Done]] = Sink.foreach(println)
source.via(flow).runWith(sink)
```

*Note.* In this example, **mapAsync** allows for asynchronous processing of each element in the stream, maintaining responsiveness even with time-consuming operations. **Reactive Streams** are designed to **handle data asynchronously**, allowing for non-blocking operations. This is crucial for building responsive systems that can handle high volumes of data without becoming unresponsive. *Creator:* Author's own work.

Table 3.14: Non-blocking Back Pressure

```
import akka.stream.OverflowStrategy
val bufferingFlow: Flow[Int, Int, NotUsed] = Flow[Int]
  .buffer(10, OverflowStrategy.backpressure)
  .map { x =>
    Thread.sleep(100) // Simulate slow processing
    x * 2}
Source(1 to 100)
  .via(bufferingFlow)
  .runWith(Sink.foreach(println)))
```

*Note.* In this example, the buffer stage with `OverflowStrategy.backpressure` ensures that if the downstream processing is slow, the upstream will be notified to slow down its production rate. **Back pressure** is a mechanism that allows consumers to signal to producers about their current capacity to handle data. This prevents fast producers from overwhelming slow consumers, ensuring system stability. *Creator:* Author's own work.

<sup>33</sup>Kuhn, R., Hanafee, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications

Table 3.15: Stream Composability

```
def filterEven: Flow[Int, Int, NotUsed] = Flow[Int].filter(_ % 2 == 0)
def multiply(factor: Int): Flow[Int, Int, NotUsed] = Flow[Int].map(_ * factor)
def logElement: Flow[Int, Int, NotUsed] = Flow[Int].map { x => println(s"Processing: $x") x}
val composedFlow: Flow[Int, Int, NotUsed] = filterEven.via(multiply(2)).via(logElement)
Source(1 to 10).via(composedFlow).runWith(Sink.foreach(println))
```

*Note.* This example demonstrates how multiple simple flows can be composed to create a more complex data processing pipeline.

**Reactive Streams** are designed to be composable, allowing complex data processing pipelines to be built from simpler, reusable components. *Creator:* Author's own work.

Table 3.16: Bounded Resource Consumption

```
import akka.stream.ThrottleMode
val throttledSource: Source[Int, NotUsed] = Source(1 to 100).throttle(10, 1.second, 5, ThrottleMode.shaping
)
throttledSource.runWith(Sink.foreach(println))
```

*Note.* This example uses throttling to limit the rate of data flow, ensuring bounded resource consumption even with fast producers. **Reactive Streams** ensure that system resources are used efficiently by **bounding** the number of elements that can be in-flight at any given time. This prevents memory overflow and system crashes due to unbounded queues. *Creator:* Author's own work.

Table 3.17: Bounded Resource Consumption

```
import akka.stream.scaladsl._
import akka.{Done, NotUsed}
import scala.concurrent.Future
case class SensorData(id: String, value: Double)
val sensorSource: Source[SensorData, NotUsed] = Source.repeat(0).map { i =>
  SensorData(s"sensor-${i % 10}", math.random())
}.throttle(100, 1.second)
val aggregationFlow: Flow[SensorData, (String, Double), NotUsed] =
  Flow[SensorData]
    .groupBy(10, _.id)
    .fold((0.0, 0)) { case ((sum, count), data) => (sum + data.value, count + 1) }
    .map { case (id, (sum, count)) => (id, sum / count) }
    .mergeSubstreams
val persistSink: Sink[(String, Double), Future[Done]] =
  Sink.foreach { case (id, avg) => println(s"Sensor $id average: $avg") }
sensorSource
  .via(aggregationFlow)
  .runWith(persistSink)
```

*Note.* This example demonstrates a more complex data engineering scenario using Reactive Streams principles in Scala. It simulates sensor data processing, including throttling, grouping, aggregation, and persistence. *Creator:* Author's own work.

### 3.7 Apache Spark

Scalable and robust distributed computing was the driving force for the creation of the **Apache Sparks** framework. Spark works essentially with a master worker configuration, in which a cluster of executor nodes is divided up into work by a **central coordinator (called the driver software)**. Using the combined computing power of several computers, this method enables Spark to manage datasets efficiently (Sree Sandhya Kona, 2023)<sup>16</sup>.

**Spark** applications containing the function and taking on the duty of initializing the **SparkContext**, which functions as a gateway to all of Spark's capabilities, begins with the **driver program**. For resource and task scheduling throughout the cluster, the **SparkContext** works with cluster managers such as **YARN**, **Mesos**, or **Spark's standalone cluster manager**. Spark runs well on cluster managers thanks to this degree of abstraction; application code is not needed to be changed (Sree Sandhya Kona, 2023)<sup>16</sup>.

Operation of a Spark cluster involves **executor nodes**. Every executor runs on a worker node and is a Java Virtual Machine (JVM) process. The driver software assigns tasks to these executors, which also have to handle data storage in memory or on disk. Caching data in memory is what sets Spark apart from disk-based systems and improves algorithm and interactive data analysis performance (Sree Sandhya Kona, 2023)<sup>16</sup>.

**Resilient Distributed Datasets (RDDs)** are the foundation of **Spark's execution model**. **RDDs** are collections of records that are immutable, partitioned and can be operated on in parallel. They are the foundation of Spark's technique for fault tolerance. If a node fails and a portion of an RDD disappears, Spark may rebuild it using data about its past kept in the RDD. This approach spares Spark from costly data duplication and allows it to recover from errors (Sree Sandhya Kona, 2023)<sup>16</sup>.

A **Directed Acyclic Graph (DAG) scheduler** is used by the Spark engine to improve the execution of complicated processes. When an operation is performed on an RDD, Spark inspects its history of actions. Forms an execution plan. Substages of this plan are then made up of jobs that may execute simultaneously across the cluster. The DAG scheduler excels at optimizing operations flow and reducing data movement over the network, which significantly boosts Spark's performance edge (Sree Sandhya Kona, 2023)<sup>16</sup>.

**Spark** integrates libraries that enhance its capabilities beyond basic tasks, in addition to the essential components that make up the Spark framework. For structured data processing, **Spark SQL**; for real-time data streams, **Spark Streaming**; for machine learning tasks, **MLlib**; and for graph computations, **GraphX**. Because these libraries interface directly with the main engine, developers may combine many data processing techniques into a single application. A memory management system included into Sparks design additionally cleverly decides how to allocate memory for computational and data storage activities (Sree Sandhya Kona, 2023)<sup>16</sup>.

By use of this approach, known as **Unified Memory Management**, Spark may modify the memory allocation

---

<sup>16</sup>Sree Sandhya Kona. (2023). Leveraging spark and pyspark for data-driven success: Insights and best practices including parallel processing, data partitioning, and fault tolerance mechanisms. *Journal of Mathematical & Computer Applications*. <https://api.semanticscholar.org/CorpusID:270301015>

for execution and caching according to the workload. Spark's performance in managing workloads ranging from batch processing to interactive queries is partly attributed to this adaptive memory management (Sree Sandhya Kona, 2023)<sup>16</sup>.

Spark's modular architecture also translates to its data source and format adaptability. Among the many data sources Spark can read from and write to are cloud storage systems, conventional relational databases and distributed file systems like HDFS. This adaptability makes Spark possible to be integrated into current data contexts, makes complex data processing pipelines that straddle many storage systems and data types possible (Sree Sandhya Kona, 2023)<sup>16</sup>.

### 3.7.1 Spark DataFrames and APIs

With its higher-level abstraction approach than the fundamental Resilient Distributed Datasets (RDDs), Spark **DataFrames** and **Datasets APIs** represent a breakthrough in Spark data processing capabilities. Using structured and semi-structured data should be made easy for users with the help of these APIs. They combine working with local data sets with the strength of Spark distributed computation (Chambers and Zaharia, 2018)<sup>35</sup>.

Introduced in Spark 1.3 **DataFrames** provide a distributed collection of data structured into named columns. Essentially, they are **equivalent to tables in a database or data frames in R/Python**. Enhanced by Spark distributed computing prowess. DataFrames are tailored for handling of large-scale structured data. They come with a domain-specific language (**DSL**) for manipulating distributed data allowing operations, like filtering, grouping and aggregation to be expressed naturally and concisely compared to RDDs (Chambers and Zaharia, 2018)<sup>35</sup>.

Using optimization techniques to produce effective query plans, Spark SQL's Catalyst optimizer is used by the **DataFrame API**. Behind-the-scenes optimization allows developers to concentrate on communicating their data modifications while Spark manages the effective cluster-wide execution. Predicate pushdown, column trimming and reordering joins are just a few of the improvements the Catalyst optimizer can undertake to significantly improve query speed in data processing processes (Chambers and Zaharia, 2018)<sup>35</sup>.

Introduced with Spark 1.6, **datasets** extend the DataFrame notion by providing a strongly-typed API that combines the benefits of RDDs type safety with the DataFrame API's efficiency. Because datasets make advantage of Scala compile-time type safety, they are especially well-suited for usage with Scala. This feature produces dependable code by allowing problems to be found at compile time that would only be detected at runtime with DataFrames (Chambers and Zaharia, 2018)<sup>35</sup>.

A DataFrame in Scala is just an untyped JVM object represented by Row in a type alias for Dataset[Row].

---

<sup>16</sup>Sree Sandhya Kona. (2023). Leveraging spark and pyspark for data-driven success: Insights and best practices including parallel processing, data partitioning, and fault tolerance mechanisms. *Journal of Mathematical & Computer Applications*. <https://api.semanticscholar.org/CorpusID:270301015>

<sup>35</sup>Chambers, B., & Zaharia, M. (2018). Spark: The definitive guide: Big data processing made simple. <https://api.semanticscholar.org/CorpusID:59220015>

Scala's DataFrames and Datasets integration facilitates a switch between typed and untyped APIs, letting developers choose the degree of type safety that most closely fits their needs. Compile time type safety and the Catalyst optimizer's improvements may be benefited by developers working with datasets that include case classes (Chambers and Zaharia, 2018)<sup>35</sup>.

DataFrames and Datasets have the same crucial feature of being able to automatically infer schemas from data sources. Working with data sources is made easier by the inferred schemas and Spark SQL's support for reading data formats such as JSON, Parquet and Avro. When automated schema inference fails, developers may manually construct schemas using both APIs to have control over column names and data types (Chambers and Zaharia, 2018)<sup>35</sup>.

The DataFrames and Datasets APIs also provide a variety of data manipulation and analysis capabilities. Along with more sophisticated chores like window functions and intricate type manipulations, these operations encompass filtering, sorting and aggregating data. Many of these procedures are stated in easily understandable domain-specific languages (DSLs). This opens up complicated data transformations to developers without distributed computing experience (Chambers and Zaharia, 2018)<sup>35</sup>.

Moreover, interoperability of the DataFrames and Datasets APIs with Spark components is one of their main benefits. Building machine learning pipelines may be streamlined, for instance, by the smooth interaction of machine learning algorithms in Spark MLlib with DataFrames. In a similar vein, batch and stream processing may be integrated inside of a single application using Spark Streaming's ability to generate and consume data frames (Chambers and Zaharia, 2018)<sup>35</sup>.

When managing data processing chores, DataFrames and Datasets often perform better than RDDs. Their power to use the improved execution engine of Spark SQL, which includes rule- and cost-based improvements, is what makes them successful. Datasets' encoders also allow JVM objects to be serialized and deserialized, which lowers memory use and improves speed when managing complex data formats (Chambers and Zaharia, 2018)<sup>35</sup>.

---

<sup>35</sup>Chambers, B., & Zaharia, M. (2018). Spark: The definitive guide: Big data processing made simple. <https://api.semanticscholar.org/CorpusID:59220015>

## Chapter 4

# Stream processing with Scala

Organizations are now able to extract value from high-velocity, continuous data flows in real-time or near real-time thanks to the concept of stream processing, which has emerged as a crucial paradigm in contemporary data engineering. This chapter explores the fundamental concepts, evolution and practical applications of stream processing.

This chapter starts with an examination of the fundamental concepts of stream processing. These principles include the nature of data streams, the difference between event time and processing time, windowing strategies, state management and fault tolerance mechanisms. These ideas are the building blocks for knowing how stream processing systems work and the problems they try to solve.

The chapter then charts the development of data processing paradigms, starting with conventional batch processing and ending with micro-batch processing and the real-time features provided by contemporary stream processing frameworks. Within the context of today's fast-paced corporate settings, this historical backdrop serves to show the rising demand for instantaneous data analysis and decision-making.

It next looks at the important function stream processing plays in modern big data ecosystems and data engineering. Managing high-velocity data and facilitating event-driven architectures, stream processing has evolved into a vital tool for companies handling massive, time-sensitive data.

The chapter also explores the difficulties and factors to be taken into account when putting stream processing systems into practice, like managing out-of-order or late data, striking a balance between throughput and latency and making sure fault tolerance in distributed settings.

Lastly, go over well-known stream processing frameworks and technologies, focusing especially on those that use the Scala programming language. This is a summary of the main capabilities and applications of systems like Apache Spark Structured Streaming, Akka Streams and Apache Flink.



## 4.1 Definition and Concepts

The concept of stream processing is intended to manage continuous data flows in real or near real time. Processed in big batches at predetermined times, it entails assessing and acting on data as it comes in. It helps companies to move quickly and get quick insights from data. Situations calling for quick data analysis and decision-making find it very helpful.

1. **Data Streams** — Continuous, unbounded sequences of data that are generated and processed in real time.
2. **Event Time vs. Processing Time** — Distinguishing between when an event occurred and when it's processed, which is crucial for handling out-of-order events.
3. **Windowing** — Grouping data into finite time-based or count-based windows for analysis.
4. **State Management** — Maintaining and updating state information as new data arrives.
5. **Fault Tolerance** — Ensuring data processing continuity and correctness in the face of failures.
6. **Continuous computation** — Ongoing analysis of unbounded data streams
7. **Real time processing** — Analyzing data immediately as it arrives.
8. **Low latency** — Minimizing delay between data ingestion and results.
9. **Stateful computations** — Maintaining state across streaming events.

## 4.2 Introduction to Stream Processing

### 4.2.1 Evolution of Data Processing Paradigms

Stream processing evolution can be broadly categorized into three main stages: batch processing, micro-batch processing and stream processing.

For many years, managing data was done best via batch processing. It required gathering data over time and processing it in batches at prearranged intervals. With big amounts of data that didn't need instantaneous results, this method performed well. Jobs requiring efficient batch processing of information included weekly sales reports and end-of-day financial reconciliations. The drawback, meantime, was the waits between data generation and insight acquisition, which may be problematic in the hectic corporate environment of today (Akidau et al., 2015)<sup>36</sup>.

---

<sup>36</sup>Akidau, T., Bradshaw, R. W., Chambers, C., Chernyak, S., Fernández-Moctezuma, R., Lax, R., McVeety, S., Mills, D., Perry, F., Schmidt, E., & Whittle, S. (2015). The dataflow model: A practical approach to balancing correctness, latency, and cost in massive-scale, unbounded, out-of-order data processing. *Proc. VLDB Endow.*, 8, 1792–1803. <https://api.semanticscholar.org/CorpusID:17844806>

A compromise between batch and real-time processing, micro batch processing emerged to meet the need for insights. In micro-batch processing, data is gathered in tiny batches and processed frequently—typically every few seconds to minutes. While keeping some of the group data processing efficiency features, this approach lowers latency as compared to batch processing. Offering insights appropriate for applications needing information in almost real time with acceptable delays, micro-batch processing achieves a balance (Akidau et al., 2015)<sup>36</sup>.

Real time or near real time data analysis is made possible by stream processing, the most recent development in data handling techniques. Because data is analyzed as soon as it comes in stream processing, insights and actions may be taken instantly. When identifying fraud, offering real-time recommendations, or monitoring systems, this approach works effectively. Effective management of high-velocity data streams has been made feasible by stream processing solutions like Apache Kafka, Apache Flink and Apache Spark Streaming, allowing companies to react to events as they happen (Akidau et al., 2015)<sup>36</sup>.

Every one of these techniques is used in contemporary data architectures. Selecting one of them over the other is contingent upon the particular requirements of the use case. Mixing these techniques, many companies utilize stream processing for insights and actions, micro-batch for reporting in nearly real time and batch processing for historical analysis (Akidau et al., 2015)<sup>36</sup>.

Growing significance of real-time data in corporate decision-making and advancements in hardware capabilities in distributed computing technologies have driven the development of these techniques. We can anticipate developments in data processing as long as data quantities keep increasing and the need for insights keeps increasing. Maybe as a result, these strategies will be combined. Provide more flexible and effective methods of handling large amounts of data (Akidau et al., 2015)<sup>36</sup>.

### 4.3 Stream Processing in Modern Data Engineering

Real-time insights and faster decision-making are two of stream processing's main benefits. In many current corporate contexts, data value dramatically declines with time. Instead of waiting for batch processing operations, stream processing enables businesses to examine and act on data in real time. Particularly crucial for real-time fraud detection in the financial industry, customized suggestions in e-commerce, predictive maintenance in manufacturing and ongoing patient monitoring in healthcare, is this capability. Businesses have a significant competitive advantage when they can see and deal with events, trends, or anomalies as they happen thanks to real-time data processing (Akidau et al., 2018)<sup>37</sup>.

Managing the high-velocity data that defines many current data sources also requires stream processing. Data is being produced at previously unheard-of rates with the ubiquity of social media, financial market data, IoT devices and online clickstreams. This amount and speed of data are often too much for conventional batch processing

---

<sup>37</sup>Akidau, T., Chernyak, S., & Lax, R. (2018). *Streaming Systems The What, Where, When, and How of Large-Scale Data Processing*. O'Reilly Media

systems to handle. Because stream processing frameworks are designed to ingest and analyze millions of events per second, businesses can extract value in real time and keep up with the rate of data creation. Situations where this skill is especially crucial include tracking user activity on websites, social media trends analysis and industrial sensor monitoring (Akidau et al., 2018)<sup>37</sup>.

Significantly less delay between data production and insight derivation is another major benefit of stream processing. In many applications, like real-time advertising bidding, instantaneous fraud detection in financial transactions, or quick reaction to user inputs in mobile apps, even little delays may have serious repercussions. Stream processing cuts down on this delay by handling data as it comes in. This makes businesses more flexible and able to adapt to new situations and customer needs. Better customer experiences, higher operational effectiveness and new company prospects may all result from this greater responsiveness (Akidau et al., 2018)<sup>37</sup>.

Comparing stream processing to batch processing, the former also helps to use resources more effectively. Stream processing distributes the computing burden over time rather than doing big, resource-intensive batch processes at predetermined intervals. More constant and predictable resource use results from this strategy, which may save money, particularly in cloud settings where resources are charged according to utilization. Furthermore, since less data has to be stored for subsequent processing, the capacity to process data incrementally as it enters often leads to decreased end-to-end latency and storage needs (Akidau et al., 2018)<sup>37</sup>.

Further supporting event-driven architectures, which are becoming more and more common in modern software systems, is stream processing. Event-driven architectures provide scalability, quick reactions to real-time events and flexible linkages between components. The foundation of these designs' management of the continuous stream of events is stream processing, which enables systems to quickly adjust to changes in information or commercial conditions. Particularly useful is this feature in microservices systems, where many services need to communicate and react quickly to events in real time (Akidau et al., 2018)<sup>37</sup>.

The latest data integration scenarios also heavily rely on stream processing. With data produced in several forms and dispersed throughout many systems, stream processing may act as a unifying layer. Data from several sources may be ingested, transformed on the fly and sent in the necessary format to different locations. Consistency across systems, real-time dashboards and prompt data-driven choices throughout the company depend on this real-time data integration capacity (Akidau et al., 2018)<sup>37</sup>.

Stream processing is becoming more and more significant in the framework of machine learning and artificial intelligence. Online learning is made possible by it and models may be modified instantly according to new data. When circumstances are changing quickly, like in the case of fraud detection or recommendation systems, this capacity is essential. Stream processing also makes machine learning models score in real time, enabling the instantaneous application of AI insights to arriving data streams (Akidau et al., 2018)<sup>37</sup>.

Finally, dealing with the "long tail" of data processing requirements is impossible without stream processing. Despite the fact that batch processing is still useful for a variety of situations, there are a great number of use cases in which the capability to process and respond to data in real time is absolutely necessary. This gap is filled by

stream processing, a versatile and potent tool for managing a broad spectrum of data processing requirements that don't fit well into the batch processing paradigm (Akidau et al., 2018)<sup>37</sup>.

## 4.4 Stream Processing in Big Data Ecosystems

Stream processing (SP) frameworks are intended to handle high-velocity, continuous data streams from several sources like log records, social media feeds, financial transactions and Internet of Things devices in the context of big data ecosystems. These frameworks allow enterprises to quickly extract value from their data resources by promptly ingesting, processing and analyzing data (Wingerath et al., 2016)<sup>29</sup>.

The ability of SP to reduce latency between data production and insight extraction is one important advantage of using it in big data settings. Because data is gathered, stored and then analyzed in big chunks, traditional batch processing methods sometimes result in significant delays. By contrast, SP acts on arriving data right away, enabling real-time analytics and quick decision-making. Particularly important in situations where quick action may provide significant commercial results include fraud detection, rapid advice and predictive maintenance (Wingerath et al., 2016)<sup>29</sup>.

The main goals of SP frameworks in big data environments are often fault tolerance and scalability. Through the distribution of processing across machine clusters, they can handle large volumes of data and provide effective throughput rates with few delays. Often included in these frameworks are features like automatic load balancing, data partitioning and fault recovery techniques to maintain system performance and dependability in the event of hardware malfunctions or network problems (Wingerath et al., 2016)<sup>29</sup>.

Key to SP is integration with other components of the big data ecosystem. These days' SP frameworks are designed to integrate with a variety of message queues, analytics tools and data storage systems with ease. They may, for example, analyze data in real time from distributed messaging systems like Apache Kafka and store the results in distributed databases or data lakes for further analysis (Wingerath et al., 2016)<sup>29</sup>.

To further ease the creation of streaming applications, several SP frameworks additionally include DSLs (Domain-Specific Languages) and user-friendly APIs. Frequently, these APIs have pre-built features for typical SP chores as windowing, merging, filtering and aggregation. Developers can focus on the core logic of their streaming apps instead of getting bogged down in the details of distributed systems when they use this level of abstraction (Wingerath et al., 2016)<sup>29</sup>.

Another important thing about SP in big data environments is that it can process both real-time and past data using the same structure. This method, called "lambda architecture" or "kappa architecture," lets businesses get the

---

<sup>37</sup>Akidau, T., Chernyak, S., & Lax, R. (2018). *Streaming Systems The What, Where, When, and How of Large-Scale Data Processing*. O'Reilly Media

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>

best of both worlds: the speed and accuracy of real-time processing and the thoroughness and accuracy of group processing (Wingerath et al., 2016)<sup>29</sup>.

Within large data settings, machine learning and AI integration is becoming more and more important in SP. A lot of streaming services nowadays provide instantaneous deployment and modification of machine learning models, enabling continuous learning and modification using real-time data. Advanced uses including real-time anomaly detection, predictive analysis and tailored recommendations are made possible by this (Wingerath et al., 2016)<sup>29</sup>.

We're witnessing tendencies toward even reduced latency, better throughput and more advanced processing capabilities as SP develops. With the integration of ideas like edge computing, network latency and bandwidth consumption may be lowered by processing data closer to the source (Wingerath et al., 2016)<sup>29</sup>.

## 4.5 Challenges and Considerations in Stream Processing

Controlling data that comes late or out of order is one of the challenges in stream processing. Data might get to nodes in distributed systems at various times because of system failures, network delays, or variable processing rates. If this case isn't treated properly, it could lead to undesirable outcomes. Time windowing and watermarking techniques must be included into stream processing systems in order to overcome this problem. Event grouping according to their arrival or occurrence time is made possible by time frames. Conversely, watermarks provide a way to gauge how far behind the system may be in processing events, which makes managing late data more precisely possible (Ounacer et al., 2017)<sup>42</sup>.

A major additional difficulty is keeping distributed stream processing nodes in a consistent state. Whereas state is usually reset between task runs in batch processing, stream processing often has to preserve long-lived state across continuous data flows. In distributed systems, where many nodes may be processing various stream segments at the same time, this becomes quite complicated. When a node fails or a network divides, methods like distributed state storage, check pointing and precisely once processing semantics are essential to guaranteeing correct and consistent results (Fragkoulis et al., 2023)<sup>43</sup>.

In stream processing, scalability and fault tolerance bring up even another set of factors. The system has to be able to grow horizontally by adding additional processing nodes as data quantities and velocities rise. For equal workload distribution across the cluster, meticulous design of data partitioning techniques is necessary. Furthermore, the system has to be resistant to network problems, node failures and other typical distributed environment disturbances. Building strong stream processing pipelines requires putting fault-tolerant techniques like data repli-

---

<sup>29</sup>Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>

<sup>42</sup>Ounacer, S., Talhaoui, M. A., Ardchir, S., Daif, A., & Azouazi, M. (2017). Real-time Data Stream Processing Challenges and Perspectives. *IJCSI International Journal of Computer Science Issues*, 14(5)

cation, automated failover and state recovery into place (Fragkoulis et al., 2023)<sup>43</sup>.

In stream processing systems, accuracy, throughput and latency are all constantly balancing challenges. Real time applications frequently need low latency, but getting there might sometimes mean sacrificing accuracy or performance. For example, although processing data in smaller batches may lower latency, it can also lower throughput overall. Analogously, approximation algorithms may provide speedier answers at the expense of some accuracy. According to the particular needs of their use cases, data engineers must carefully tune their systems, often compromising between these conflicting elements (Fragkoulis et al., 2023)<sup>43</sup>.

In stream processing systems, accuracy, throughput and latency are all constantly balancing challenges. Real time applications frequently need low latency, but getting there might sometimes mean sacrificing accuracy or performance. For example, although processing data in smaller batches may lower latency, it can also lower throughput overall. Analogously, approximation algorithms may provide speedier answers at the expense of some accuracy. Data engineers are required to carefully tune their systems in accordance with the particular needs of their use cases and they often have to make compromises between the many aspects that are fighting for their attention (Fragkoulis et al., 2023)<sup>43</sup>.

Long-running stream processing programs are greatly challenged by changing schemas and data models. The format and semantics of the streaming data might vary as well as business needs do. System changes must be accommodated by stream processing systems without needing system downtime or losing past data. Usually, this means putting data schema versioning systems into place and making sure data processing logic is backward compatible (Fragkoulis et al., 2023)<sup>43</sup>.

The distributed and real-time character of stream processing programs may make debugging and monitoring them especially difficult. In streaming settings, conventional batch processing debugging methods might not be useful or effective. Putting in place thorough monitoring, logging and metrics collecting systems is essential to finding and fixing problems in production settings. Debugging and testing stream processing algorithms may also benefit much from tools for replaying old data streams (Fragkoulis et al., 2023)<sup>43</sup>.

Particularly, when handling regulated data, stream processing systems must take security and compliance into account. Use of audit trails, access limits and data encryption is essential. When data is constantly moving and being processed in real time, as in streaming scenarios, it might be difficult to comply with data protection regulations like GDPR or CCPA (Ounacer et al., 2017)<sup>42</sup>.

Moreover, stream processing system operation might be somewhat difficult. For stream processing, this often requires unique frameworks, real-time processing methods and specialist understanding networked systems. To get the most out of stream processing technologies, businesses should spend money on training and hiring experts (Fragkoulis et al., 2023)<sup>43</sup>.

---

<sup>43</sup>Fragkoulis, M., Carbone, P., Kalavri, V., & Katsifodimos, A. (2023). A survey on the evolution of stream processing systems. *VLDB Journal*

<sup>42</sup>Ounacer, S., Talhaoui, M. A., Ardchir, S., Daif, A., & Azouazi, M. (2017). Real-time Data Stream Processing Challenges and Perspectives. *IJCSI International Journal of Computer Science Issues*, 14(5)

## 4.6 Stream Processing Technologies and Frameworks

### Popular Stream Processing Frameworks

- Apache Kafka Streams: Lightweight library for building streaming applications
- Apache Flink: Distributed processing engine for stateful computations over data streams
- Apache Spark Structured Streaming: Stream processing on top of the Spark engine
- Apache Storm: Distributed real time computation system

### SQL for Stream Processing

- KSQL for Kafka Streams
- Flink SQL
- Spark Structured Streaming SQL

### Cloud-based Streaming Services

- Amazon Kinesis
- Google Cloud Dataflow
- Azure Stream Analytics

## 4.7 Popular Stream Processing Frameworks in Scala

**Apache Spark Structured Streaming** is a scalable fault-tolerant stream processing engine that is constructed on the foundation of the Spark SQL engine. It offers a user DataFrame-based API that lets developers express streaming computations the same way they would express batch computations. **Key features:** Unified API for batch and streaming; End-to-end exactly once guarantees; Event-time processing; Incremental query planning and execution.

**Akka Streams** is a library built on top of Akka actors, providing a higher-level abstraction for building streaming applications. It implements the Reactive Streams specification, offering back-pressure and asynchronous processing. **Key features:** Composable stream processing DSL; Integration with Akka actors; Built-in back-pressure support; Rich set of stream operators.

**Apache Flink** is a distributed stream processing framework that provides low-latency, high-throughput data processing. While primarily developed in Java, it offers a Scala API that leverages Scala's language features. **Key features:** Event-time processing and watermarks; Stateful computation; Exactly-once processing semantics; Savepoints for application versioning.



## Chapter 5

# Testing Data Pipelines

Complex procedures of extracting, converting and loading enormous amounts of data from many sources are part of data engineering. Inadequate testing exposes these procedures to mistakes that may spread across the whole data ecosystem, resulting in inaccurate analysis and poor decision-making. The expense and work of resolving difficulties in production settings are reduced when testing helps find and repair problems early in the development cycle (Punn et al., 2019)<sup>38</sup>.

Testing becomes much more important for Scala data pipelines built using functional programming. Testable code is made possible in functional programming by the usage of pure functions and the immutability of data. Data transformations are consistent and repeatable when testing confirms that these functional concepts are implemented appropriately. In addition, property-based testing—which works well with functional programming paradigms—can reveal edge situations and unanticipated actions in data processing algorithms (Punn et al., 2019)<sup>38</sup>.

In data engineering, comprehensive testing also tackles issues with data quality. All along the pipeline, it makes sure that data follows anticipated formats, ranges and business requirements. In big data situations, where the volume and variety of data might make human verification impossible, this is especially crucial. Continuously assessing the quality of the data, automated testing may send out notifications when abnormalities are found (Punn et al., 2019)<sup>38</sup>.

Moreover, preserving the scalability and performance of data engineering solutions depends critically on testing. While scalability testing guarantee that systems can manage growing data quantities, performance tests may find data processing bottlenecks. Testing in the setting of Scala and functional programming may confirm that the system makes effective use of resources and that parallel processing capabilities are applied effectively (Punn et al., 2019)<sup>38</sup>.

---

<sup>38</sup>Punn, N. S., Agarwal, S., Syafrullah, M., & Adiyarta, K. (2019). Testing Big Data Applications. *Proceeding of the Electrical Engineering Computer Science and Informatics*

## 5.1 Testing Levels for Data Pipelines

Data pipeline testing pyramid starts with unit testing. This level tests each function and transformation separately. Because Scala emphasizes pure functions and immutability, it works very well for unit testing. Little, independent logic parts may be effortlessly tested by developers without concern about side effects or state changes. This enables complete validation of the atomic functions that comprise the pipeline, including filtering operations and data transformation algorithms. Unit tests may be written in Scala using its robust testing frameworks, such as `ScalaTest` (O’Keeffe, 2021)<sup>44</sup>.

As one advances through the testing hierarchy, integration testing is concerned with confirming the connections between various data pipeline components. This degree of testing guarantees that transformations function as intended when coupled and that data moves appropriately across pipeline stages. Integration tests in the setting of Scala-based data engineering might include confirming the composition of many functional transformations or examining the interactions between various pipeline modules. Catching problems that could develop from the integration of individually correct but maybe incompatible components depends on this degree of testing (O’Keeffe, 2021)<sup>44</sup>.

At the highest level, End-to-End Testing validates the functionality of the entire data pipeline from ingestion to output. To be sure the pipeline provides the desired outcomes, data must be passed through the whole pipeline and real-world situations are simulated. End-to-end tests may take use of the composability of functions to build extensive test scenarios for data pipelines built in Scala utilizing functional programming ideas. For the purpose of ensuring that the pipeline satisfies the objectives of the company and appropriately manages a variety of data circumstances, these tests are absolutely necessary (O’Keeffe, 2021)<sup>44</sup>.

Data pipelines sometimes need specific testing methodologies in addition to the conventional levels of testing. Testing data quality, for example, is concerned with confirming that data is accurate, consistent and integrity across the pipeline. Assuring that the pipeline can manage anticipated data volumes and satisfy processing time requirements requires performance testing. Particularly helpful for effectively creating and running these kinds of tests is Scala’s support for concurrent programming and parallel collections (O’Keeffe, 2021)<sup>44</sup>.

## 5.2 Testing Frameworks and Tools for Scala

**ScalaTest** is one of the most widely used and adaptable Scala testing frameworks. Supporting several testing methods — like `FunSuite`, `FlatSpec` and `WordSpec` — it offers a versatile and expressive approach to design tests. `ScalaTest`’s rich set of matchers allows for clear and readable assertions, making it easier to express complex test conditions. Because `ScalaTest` allows integration with other tools and frameworks, it is especially helpful for data engineering when testing different parts of a data pipeline.

---

<sup>44</sup>O’Keeffe, D. (2021). The test pyramid, data engineering, and you! <https://servian.dev/the-test-pyramid-and-data-engineering-with-julia-e4678c3f8dff>

Allowing property-based testing, **ScalaCheck** is another potent tool that enhances ScalaTest. When working with intricate data transformations is a common task in functional programming and data engineering, this method is quite helpful. ScalaCheck enables the definition of properties that code must satisfy and the automatic generation of test cases to verify these properties. Specifying invariants that should hold true for any input data makes this especially helpful for testing data transformations.

**Mockito** and other frameworks may be used with ScalaTest to mock external dependencies in tests. When testing data pipeline components that communicate with databases, APIs, or other systems, this is critical. Mocking lets run several scenarios and isolate the component under test without relying on actual external resources.

Performance testing of data pipelines may benefit from **Scala's integrated support for concurrency and parallel collections**. Even though these language features aren't testing frameworks themselves, they let write tests that check scalability and efficiency of the data processing logic when there are different loads conditions.

When a data pipeline includes Spark processing, tools such as the testing utilities provided by **Apache Spark** may be used for integration and end-to-end testing. The ability to create tests that mimic distributed data processing situations made possible by these technologies is essential for verifying pipeline functionality in a production-like setting.

Finally, Scala tests may be set up to run automatically by continuous integration systems like **Jenkins** or **GitLab CI**, guaranteeing the dependability of data pipeline even while making adjustments and enhancements. Using these tools unit tests, integration tests, and even performance benchmarks may be configured to run as part of development process using these tools.

## Chapter 6

# Deploying and monitoring data pipelines

A vital component of data engineering, deployment and monitoring guarantee the performance, reliability and seamless functioning of data pipelines in production settings. It is impossible to exaggerate the value of strong deployment and monitoring procedures as data pipelines become more intricate and vital to businesses. Data engineers can easily move their Scala-based data pipelines from development to production with the help of good release strategies. These strategies make sure that the pipelines can handle the numbers and speeds of real-world data while keeping the data safe and secure. Conversely, good monitoring gives insight into the behavior, performance and health of data pipelines, enabling early detection and fixing of problems before they affect company operations (Tome et al., 2024)<sup>10</sup>.

Within the field of data engineering, deployment includes not only the act of transferring code and settings to production environments but also the coordination of many elements, such as data sources, processing engines, storage systems and downstream applications. This often includes continuous integration/continuous deployment (CI/CD) pipelines, orchestration systems like Kubernetes and containerization technologies like Docker for Scala-based data pipelines. These techniques and technologies allow data engineers to handle dependencies well, get consistency across environments, and quickly iterate on pipeline enhancements (Humble and Farley, 2010)<sup>46</sup>.

Equally important is monitoring data pipelines, as it offers information on the operations of data processing procedures. This involves monitoring key performance indicators (KPIs) like error rates, processing lag, data throughput and resource use. Monitoring Scala-based pipelines often requires integrating with metric collecting systems, logging frameworks and visualization tools that can manage the distributed and asynchronous character of Scala applications. By monitoring well, data engineers may optimize resource allocation, guarantee service-level agreements (SLAs) are fulfilled and resolve problems fast when they occur (Tome et al., 2024)<sup>10</sup>.

But Scala-based data pipeline deployment and monitoring have unique difficulties of their own. Managing the

---

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

<sup>46</sup>Humble, J., & Farley, D. (2010). *Continuous delivery: Reliable software releases through build, test, and deployment automation*. <https://api.semanticscholar.org/CorpusID:111609675>

intricacy of Scala’s ecosystem and its compatibility with other big data technologies is one major problem. Strong characteristics of Scala, including type system and functional programming paradigms, may result in complex pipeline designs that are hard to implement and track without specific expertise. Managing all dependencies and setups appropriately across many environments may be very difficult (Odersky et al., 2021)<sup>2</sup>.

One further problem is that many Scala-based data processing frameworks, like Apache Spark, are distributed. In order to successfully deploy and monitor distributed systems, it is necessary to give careful thought to a number of issues, including data partitioning, cluster management and fault tolerance. Strong as they are, Scala’s concurrency models may also make monitoring and debugging more difficult, particularly when asynchronous operations and actor-based systems are involved (Jules S Damji et al., 2020)<sup>47</sup>.

Moreover, deployment and monitoring procedures need to always change due to the quick development of Scala and the larger data engineering environment. For data engineering teams, staying current with new versions of Scala, its libraries and related tools while maintaining stability and backward compatibility may be rather difficult (Tome et al., 2024)<sup>10</sup>.

Finally, it might be difficult to combine observability with speed optimization. Although Scala enables very fast and highly optimized data pipelines, careful design and implementation are necessary to instrument these pipelines for extensive monitoring without adding a lot of overhead (Tome et al., 2024)<sup>10</sup>.

## 6.1 Cloud Deployment

**Cloud deployment** is becoming an increasingly popular and necessary strategy for enterprises seeking to extend their data engineering pipelines and infrastructure. Companies may use adaptable, scalable and reasonably priced resources for building and maintaining their data pipelines by using cloud platforms such as Google Cloud Platform (GCP), Microsoft Azure, or Amazon Web Services (AWS) (Tome et al., 2024)<sup>10</sup>.

Easy resource scaling up or down according to demand is one of the main advantages of cloud deployment for data engineering. Cloud providers include managed Hadoop clusters, data warehouses and stream processing engines among other services created especially for big data processing, storage and analytics. This frees data engineers to concentrate on developing and refining their pipelines rather than overseeing the supporting infrastructure (Pondel, 2013)<sup>45</sup>.

Organizations may also put strong disaster recovery and high availability plans into practice with cloud deployment. Usually, cloud providers provide many geographical areas and availability zones, which enables data

---

<sup>2</sup>Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press

<sup>47</sup>Jules S Damji, Brooke Wenig, Tathagata Das, & Denny Lee. (2020). *Learning Spark: Lightning-fast data analytics* (2nd ed.). O’Reilly Media

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

<sup>45</sup>Pondel, M. (2013). Business intelligence as a service in a cloud environment. *Federated Conference on Computer Science and Information System*, 1269–1271

pipelines to be spread across several sites for higher fault tolerance. The replication and backup capabilities included in many cloud services also make data protection and recovery procedures easier (Tome et al., 2024)<sup>10</sup>.

In data engineering, security and compliance are major issues that cloud providers supply with a variety of tools and services to meet. This covers identity and access management, encryption both during transit and at rest, and industry standard compliance certifications. Organizations may often get better data protection with these integrated security capabilities than they could with on-premises equipment (Tome et al., 2024)<sup>10</sup>.

Easy integration of many applications and technologies within the same ecosystem is another benefit of cloud deployment. For example, a data pipeline deployed on Google Cloud Platform could seamlessly combine services like Cloud Storage for storage, Dataproc for processing, BigQuery for warehousing and Data Studio for visualization. Complex data process creation and administration may be greatly simplified by this connection (Tome et al., 2024)<sup>10</sup>.

A major advantage of cloud implementation is also cost optimization. Pay-as-you-go pricing plans and the capacity to dynamically scale resources according to consumption allow businesses to often save money on infrastructure overall as compared to maintaining on-premises systems. To assist optimize costs and find possible savings, cloud providers can provide a range of cost management services and tools (Pondel, 2013)<sup>45</sup>.

With that being said, cloud deployment is not without its own unique set of difficulties. Data engineers should be aware of any vendor lock-in and make sure that their pipelines are made portable. To prevent unanticipated expenditures, cloud cost management also calls for vigilant monitoring and optimization (Tome et al., 2024)<sup>10</sup>.

---

<sup>45</sup>Pondel, M. (2013). Business intelligence as a service in a cloud environment. *Federated Conference on Computer Science and Information System*, 1269–1271

<sup>10</sup>Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing

## Chapter 7

# Conclusion

In this thesis, the strong synergy between Scala functional programming and contemporary data engineering techniques has been investigated. Building robust, scalable and maintainable data pipelines has been shown by means of a thorough analysis of Scala's fundamental concepts, data engineering workflows and advanced processing paradigms.

The journey began with an exploration of Scala's fundamental features, including immutable data structures, higher-order functions and lazy evaluation. Functional programming is built on these notions, which offer a strong basis for the construction of efficient data processing systems. These concepts constitute the bedrock of functional programming. Then more complex subjects such type classes, algebraic data types and pattern matching were covered, demonstrating how these capabilities support expressive data modeling and transformations.

The investigation was extended to the practical aspects of data engineering, covering the entire workflow from data ingestion to serving. A number of different paradigms for processing data were investigated, such as batch processing and stream processing. Additionally, the question of how functional programming concepts might be applied to each step of the data pipeline was investigated. Scala's potential in managing large-scale data operations were emphasized by the topic of parallel and distributed processing, notably via frameworks like Apache Spark and Akka.

A significant portion of the work was dedicated to stream processing, reflecting its growing importance in modern data architectures. The evolution of data processing paradigms and the role of stream processing in big data ecosystems were analyzed, addressing challenges and considerations unique to this approach.

The critical aspects of testing, deploying and monitoring data pipelines were also emphasized. Various testing frameworks and methodologies tailored for Scala were explored, underlining the importance of robust testing practices in ensuring the reliability and correctness of data processing systems.

In conclusion, it has been demonstrated that functional programming in Scala offers a powerful toolkit for

addressing the complexities of modern data engineering. The language's blend of object-oriented and functional paradigms, coupled with its rich ecosystem of libraries and frameworks, positions it as an excellent choice for building scalable, efficient and maintainable data processing systems.

As data continues to grow in volume, velocity and variety, the principles and techniques discussed in this thesis will become increasingly valuable. A solid foundation for tackling the challenges of big data is provided by functional programming in Scala, enabling data engineers to build systems that are not only performant but also adaptable to the ever-changing landscape of data processing requirements.

Moving forward, further integration of functional programming concepts is likely to be seen in the field of data engineering, as well as the continued evolution of stream processing technologies. Future research could explore emerging paradigms in distributed computing, advancements in real-time analytics and the application of functional programming principles to new domains within data engineering.



# Bibliography

- Achanta, A., & Boina, R. (2023). Data governance and quality management in data engineering. *International Journal of Computer Trends and Technology*. <https://api.semanticscholar.org/CorpusID:265699088>
- Agarwal, Y. (2023). Scala vector. Retrieved June 13, 2024, from <https://www.scaler.com/topics/scala/scalavector/>
- Akida, T., Bradshaw, R. W., Chambers, C., Chernyak, S., Fernández-Moctezuma, R., Lax, R., McVeety, S., Mills, D., Perry, F., Schmidt, E., & Whittle, S. (2015). The dataflow model: A practical approach to balancing correctness, latency, and cost in massive-scale, unbounded, out-of-order data processing. *Proc. VLDB Endow.*, 8, 1792–1803. <https://api.semanticscholar.org/CorpusID:17844806>
- Akida, T., Chernyak, S., & Lax, R. (2018). *Streaming Systems The What, Where, When, and How of Large-Scale Data Processing*. O'Reilly Media.
- Andrew Black & Nederpelt, P. van. (2020). Dimensions of Data Quality (DDQ). *DAMA NL Foundation*.
- Cédola, A. P., Rossini, R., Bosi, I., & Conzon, D. (2021). Feature engineering and machine learning modelling for predictive maintenance based on production and stop events. <https://api.semanticscholar.org/CorpusID:244727976>
- Chambers, B., & Zaharia, M. (2018). Spark: The definitive guide: Big data processing made simple. <https://api.semanticscholar.org/CorpusID:59220015>
- Chen, G., Chen, K., Jiang, D., Ooi, B. C., Shi, L., Vo, H. T., & Wu, S. (2011). E3 : An Elastic Execution Engine for Scalable Data Processing. *Journal of Information Processing*, 20(No.1), 65–76.
- Contributors. (2024). Task support. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/configuration.html#task-support>
- Fadiya, S., & Sari, A. (2018). The importance of big data technology. *International Journal of Engineering & Technology*. <https://api.semanticscholar.org/CorpusID:199019074>
- Fragkoulis, M., Carbone, P., Kalavri, V., & Katsifodimos, A. (2023). A survey on the evolution of stream processing systems. *VLDB Journal*.
- Ghosh, D. (2011). *DSLs in Action*. Manning Publications.
- Haller, P., Prokopec, A., Miller, H., Klang, V., Kuhn, R., Jovanovic, V., & Contributors. (2024). Futures and promises. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/core/futures.html>
- Hughes, J. (1990). Why Functional Programming Matters. “*Research Topics in Functional Programming*” ed. D. Turner, Addison-Wesley, 17–42.
- Humble, J., & Farley, D. (2010). *Continuous delivery: Reliable software releases through build, test, and deployment automation*. <https://api.semanticscholar.org/CorpusID:111609675>

- InterviewBit. (2022). Big data architecture – detailed explanation. <https://www.interviewbit.com/blog/big-data-architecture/>
- Jain, S. (2020). The a to z of microservice architecture. <https://www.systango.com/blog/a-z-microservice-architecture>
- Jules S Damji, Brooke Wenig, Tathagata Das, & Denny Lee. (2020). *Learning Spark: Lightning-fast data analytics* (2nd ed.). O'Reilly Media.
- Kuhn, R., Hanafee, B., & Allen, J. (2017). *Reactive Design Patterns*. Manning Publications.
- Kunasaikaran, J., & Iqbal, A. (2016). A brief overview of functional programming languages. *electronic Journal of Computer Science and Information Technology (eJCSIT)*, 6(No. 1). <https://api.semanticscholar.org/CorpusID:195989733>
- Lightbend. (2024). Streams. Retrieved June 21, 2024, from <https://www.reactivemanifesto.org/>
- Milewski, B. (2013). Functional data structures and concurrency in c++. Retrieved June 13, 2024, from <https://bartoszmilewski.com/2013/12/10/functional-data-structures-and-concurrency-in-c/>
- Nambiar, A. M., & Mundra, D. (2022). An overview of data warehouse and data lake in modern enterprise data management. *Big Data Cogn. Comput.*, 6, 132. <https://api.semanticscholar.org/CorpusID:253428554>
- Natesan, P., E, S. V., Mathivanan, S. K., Venkatesan, M., Jayagopal, P., & Allayear, S. M. (2023). A distributed framework for predictive analytics using big data and mapreduce parallel programming. *Mathematical Problems in Engineering*. <https://api.semanticscholar.org/CorpusID:256524905>
- Nayak, A. (2013). Type of nosql databases and its comparison with relational databases. <https://api.semanticscholar.org/CorpusID:15594476>
- Odersky, M., Altherr, P., Cremet, V., Dragos, I., Dubochet, G., Emir, B., McDirmid, S., Micheloud, S., Mihaylov, N., Schinz, M., Stenman, E., Spoon, L., & Zenger, M. (2006). *An Overview of the Scala Programming Language* (2nd ed.). École Polytechnique Fédérale de Lausanne (EPFL).
- Odersky, M., Spoon, L., Venners, B., & Sommers, F. (2021). *Programming in Scala* (5th ed.). Artima Press.
- O'Keeffe, D. (2021). The test pyramid, data engineering, and you! <https://servian.dev/the-test-pyramid-and-data-engineering-with-julia-e4678c3f8dff>
- Ounacer, S., Talhaoui, M. A., Ardchir, S., Daif, A., & Azouazi, M. (2017). Real-time Data Stream Processing Challenges and Perspectives. *IJCSI International Journal of Computer Science Issues*, 14(5).
- Pawar, H. (2023). Scala lazy evaluation. Retrieved June 13, 2024, from <https://www.scaler.com/topics/scala/lazy-evaluation-in-scala/>
- Pilquist, M., Bjarnason, R., & Chiusano, P. (2023). *Functional Programming in Scala* (2nd ed.). Manning publications.
- Pondel, M. (2013). Business intelligence as a service in a cloud environment. *Federated Conference on Computer Science and Information System*, 1269–1271.
- Prokopec, A., Miller, H., & Contributors. (2024). Parallel collections overview. Retrieved June 21, 2024, from <https://docs.scala-lang.org/overviews/parallel-collections/overview.html>
- Punn, N. S., Agarwal, S., Syafrullah, M., & Adiyarta, K. (2019). Testing Big Data Applications. *Proceeding of the Electrical Engineering Computer Science and Informatics*.
- Roychowdhury, S., & Sato, J. Y. (2021). Video-data pipelines for machine learning applications.

- Sree Sandhya Kona. (2023). Leveraging spark and pyspark for data-driven success: Insights and best practices including parallel processing, data partitioning, and fault tolerance mechanisms. *Journal of Mathematical & Computer Applications*. <https://api.semanticscholar.org/CorpusID:270301015>
- Stenberg, M. (2019). Functional and imperative object-oriented programming in theory and practice. <https://api.semanticscholar.org/CorpusID:197640090>
- Suereth, J. D. (2012). *Scala in Depth*. Manning Publications.
- Tome, E., Bhattacharjee, R., & Radford, D. (2024). *Data Engineering with Scala and Spark*. Packt Publishing.
- Wadler, P. (1992). Monads for functional programming.
- Wingerath, W., Gessert, F., Friedrich, S., & Ritter, N. (2016). Real-time stream processing for big data. *it - Information Technology*, 58, 186–194. <https://api.semanticscholar.org/CorpusID:30328536>
- Winn, J. (2023). Etl vs elt: Understanding the differences and making the right choice. Retrieved June 15, 2024, from <https://www.datacamp.com/blog/etl-vs-elt>
- Xie, Q., Zhang, H., Tang, Y.-r., & Lin, M. (2021). Solution ideas and practices for data governance engineering in colleges and universities. *E3S Web of Conferences*. <https://api.semanticscholar.org/CorpusID:235873740>
- Zibin, Y., Potanin, A., Ali, M., Artzi, S., Kie`zun, A., & Ernst, M. D. (2007). Object and Reference Immutability using Java Generics. *MIT Computer Science & Artificial Intelligence Lab*.

# List of Tables

1.1	Scala's transformation of operators into method calls . . . . .	8
1.2	Immutability and transforming lists . . . . .	10
1.3	Partitioning a vector based on a predicate . . . . .	11
1.4	Example of lazy evaluation and infinite sequences . . . . .	11
1.5	map . . . . .	12
1.6	flatMap . . . . .	12
1.7	filter . . . . .	12
1.8	reduce . . . . .	12
1.9	Generic data processing pipeline with higher-order functions . . . . .	13
1.10	Building a data processing DSL . . . . .	13
1.11	Lazy evaluation . . . . .	14
1.12	Infinite fibonacci stream . . . . .	14
1.13	Direct acrylic graph (DAG) . . . . .	15
1.14	Data processing workflow . . . . .	15
1.15	Data model using ADTs . . . . .	16
1.16	Pattern matching . . . . .	16
1.17	Data processing and aggregation . . . . .	17
1.18	Pattern matching and ADTs . . . . .	17

1.19	Implicit implementations	18
1.20	Generic function	18
1.21	Implicit instance	19
1.22	Monads and error handling	20
1.23	Parallel vector	21
1.24	Apache Spark data processing	21
1.25	Akka actors	22
3.1	Parallel vs. sequential processing	44
3.2	Non-associative operation	44
3.3	Non-associative operation	45
3.4	Performance tuning with Task Support	45
3.5	Futures	46
3.6	Using Async/Await	47
3.7	Error handling	47
3.8	Composition	47
3.9	Basic Actor Structure in Akka	49
3.10	Actor Hierarchies and Supervision	49
3.11	Actors Scalability and Performance	49
3.12	Actors Scalability and Performance	50
3.13	Asynchronous Processing	52
3.14	Non-blocking Back Pressure	52
3.15	Stream Composability	53
3.16	Bounded Resource Consumption	53

3.17 Bounded Resource Consumption . . . . .	53
---	----

# List of Figures

2.1	ETL and ELT . . . . .	31
3.1	Lambda Architecture . . . . .	41
3.2	Kappa Architecture . . . . .	42
3.3	Microservice Architecture Style Structure . . . . .	43
3.4	Reactive Streams . . . . .	50